



legoESM 1.0: A Differentiable Earth System Model in JAX

Technical Documentation of Equations and Implementations

legoESM Development Team

Version 1.0 — October 2026

Contents

1	Introduction	7
2	Model Architecture	8
3	Grid Systems	10
3.1	Cubed-Sphere Grid	10
3.1.1	Face-to-Cartesian mapping	11
3.1.2	Cell area	11
3.1.3	Grid spacing	11
3.1.4	Grid angle and wind rotation	11
3.1.5	Halo exchange	11
3.2	Gaussian Grid and Spectral Transforms	12
3.2.1	Grid dimensions	12
3.2.2	Associated Legendre polynomials	12
3.2.3	Spherical harmonic transforms	12
3.2.4	Spectral operators	13
3.2.5	Velocity from vorticity and divergence	13
3.3	Sigma Vertical Coordinate	13
3.3.1	Simmons–Burridge geopotential	13
3.3.2	Sigma-dot	13
3.3.3	Pressure velocity	13
3.4	Height and z^* Coordinates	14
4	Core Operators	14
4.1	Gradient	14
4.2	Divergence (flux form)	14
4.3	Vorticity	14
4.4	C-D Grid Operators (FV3-style, <code>operators_cdgrid.py</code>)	14
4.5	Laplacian and Hyperdiffusion	15
5	Finite-Volume Transport	16
5.1	MUSCL Reconstruction (2nd order, <code>halo = 1</code>)	16
5.2	PPM Reconstruction (4th order, <code>halo = 2</code>)	16
5.2.1	Fourth-order edge values	16
5.2.2	Monotonicity constraint	16
5.2.3	Colella–Woodward limiting	16

5.3	WENO5 Reconstruction (5th order, halo = 2)	16
5.3.1	Left-biased reconstruction at $i + \frac{1}{2}$	17
5.3.2	Smoothness indicators	17
5.3.3	Nonlinear weights	17
5.3.4	Right-biased reconstruction	17
5.4	Rusanov Flux	17
5.5	Well-balanced scheme	17
6	FC-Gram Spectral Operators	17
6.1	Extension operator	18
6.2	Spectral differentiation	18
6.3	C-grid-style divergence damping	18
7	Atmospheric Dynamics	18
7.1	Shallow Water Equations	18
7.2	Hydrostatic Primitive Equations (σ -coordinates)	19
7.3	Spectral Primitive Equations (Vorticity–Divergence)	19
7.4	Non-Hydrostatic Compressible Euler Equations	20
7.4.1	Split-explicit time integration	20
7.4.2	Sponge layer	20
7.4.3	Turbulence closure: CRM, LES, and DNS	21
7.5	FV3 Duo-Grid Lane	21
8	Atmospheric Physics	21
8.1	Gray Radiation (Frierson 2006)	21
8.1.1	Longwave optical depth	22
8.1.2	Two-stream LW fluxes	22
8.1.3	Shortwave (Beer–Lambert)	22
8.1.4	Heating rate	22
8.2	RRTMGP Correlated- k Radiation	22
8.3	SBM Convection (Frierson 2007)	22
8.3.1	Moist adiabatic lapse rate	23
8.3.2	Saturation mixing ratio (Tetens)	23
8.3.3	CAPE and trigger	23
8.3.4	Enthalpy-conserving relaxation	23
8.3.5	Precipitation	23
8.4	Held–Suarez Forcing	23
8.4.1	Equilibrium temperature	23
8.4.2	Newtonian relaxation	23
8.4.3	Rayleigh friction	24
8.5	Louis (1979) Turbulence	24
8.5.1	Mixing length	24
8.5.2	Richardson number	24
8.5.3	Stability functions	24
8.6	Vertical Diffusion (Implicit Thomas Algorithm)	24
8.7	Surface Fluxes	24
8.7.1	Constant coefficients (default)	25
8.7.2	Monin–Obukhov Similarity Theory (MOST)	25
8.7.3	COARE 3.0 roughness (Fairall et al. 2003)	25
8.7.4	Large & Yeager (2004) roughness	25
8.8	Additional Turbulence Schemes	26
8.8.1	Smagorinsky	26

8.8.2	YSU (Yonsei University)	26
8.8.3	CLUBB-lite	26
8.8.4	EDMF Turbulence	26
8.9	Additional Convection Schemes	26
8.9.1	Kuo (1974)	26
8.9.2	Mass-Flux Convection (Arakawa–Wu Type)	26
8.9.3	EDMF Convection	26
8.9.4	Deep Convective Adjustment (DCA)	26
8.10	Gravity Wave Drag	26
8.11	Cloud Microphysics	27
9	Ocean Dynamics	27
9.1	Boussinesq Hydrostatic Primitive Equations	27
9.1.1	Tracer equations (skew-symmetric)	27
9.1.2	Free-surface equation	27
9.2	Barotropic Solver (Forward–Backward, Semi-Implicit Coriolis)	27
9.3	Equation of State (Wright 1997)	28
9.3.1	Linear Equation of State	28
9.4	Hydrostatic Pressure and Vertical Velocity	28
9.5	Ocean Vertical Mixing	28
9.5.1	K-Profile Parameterisation (KPP, Large et al. 1994)	29
9.5.2	Richardson-Number Mixing	29
9.5.3	Constant Mixing	29
9.5.4	Prognostic TKE closures	29
9.6	Ocean Lateral Mixing	29
9.6.1	Harmonic (Laplacian) Diffusion	29
9.6.2	Biharmonic (∇^4) Diffusion	29
9.6.3	Gent–McWilliams / Redi (GM/Redi)	29
9.7	Ocean Bottom Drag	29
9.8	Ocean Surface Forcing	30
9.9	Shortwave Penetration	30
9.10	Ocean Convection	30
9.10.1	Enhanced Diffusion	30
9.10.2	Plume Model	30
9.11	Ocean Conservation Fixers	30
9.12	Tracer Advection Schemes	30
9.13	Realistic-Geometry Spinup	31
10	Time Integration	31
10.1	SSP-RK3 (Shu–Osher 1988)	31
10.2	Robert-Asselin-Williams (RAW) Time Filter (v3.7)	31
10.3	Adaptive Hyperdiffusion (v3.7)	31
10.4	Split-Explicit Methods	32
11	Parallel Runtime and Precision	32
11.1	Process and device setup	32
11.2	Halo schedules on the Voronoi (MPAS) mesh	32
11.3	Lat-lon and cubed-sphere exchanges	32
11.4	Ocean barotropic solver at scale	33
11.5	Precision policies	33

12 Land, Ice, and Coupler	33
12.1 Slab Land Model	33
12.1.1 Energy balance	33
12.1.2 Bucket hydrology	33
12.2 Sea-Ice Model	33
12.2.1 Surface temperature	33
12.2.2 Growth and melt	34
12.2.3 Concentration	34
12.2.4 EVP Rheology (Hunke & Dukowicz 1997)	34
12.2.5 Multi-category ice	34
12.3 Land Carbon Cycle and Stomatal Conductance	34
12.3.1 DALEC-990 Carbon Pools	34
12.3.2 Farquhar (1980) Photosynthesis	34
12.3.3 Stomatal conductance	34
12.4 Plant Functional Types (PFTs)	35
12.5 DifferLand Training	35
12.6 Interactive Multilayer Land on the MPAS Lane	35
12.7 Tile Coupler	35
13 Machine Learning: Spherical Fourier Neural Operator	36
13.1 Architecture	36
13.2 Spectral convolution	36
13.3 Neural and Learned Physics	37
13.3.1 NeuralPhysics	37
13.3.2 LearnedColumn	37
14 Training and Calibration	37
14.1 One model registry	37
14.2 Parameter specifications and tiers	37
14.3 No inert parameters	37
14.4 Process-level calibration	38
15 Conservation Corrections	38
15.1 Atmosphere	38
15.2 Water and Energy on the Production Lane	38
15.3 Ocean	39
16 Physical Constants	39
17 Ocean Biogeochemistry	40
17.1 Carbonate Chemistry and Air-Sea CO ₂ Exchange	40
17.2 NPZD Ecosystem Model	40
18 CMIP Infrastructure	40
18.1 CF/CMOR-Compliant Output	40
18.2 Production AMIP Configuration	41
18.3 Experiment Templates	41
18.4 Restart and Reproducibility	41
18.5 Tuning Guide	42
18.6 OMIP Simulation Pipeline	42

19 4D-Var Data Assimilation (v3.9)	42
19.1 Cost Function	42
19.2 Control Vector	42
19.3 Background Error Covariance	43
19.4 Observation Operators	43
19.5 Minimiser	43
19.6 DA Cycling and Incremental Analysis	43
20 References	43

Introduction

legoESM is a modular Earth System Model written entirely in JAX, supporting automatic differentiation through all components. The model couples atmospheric, oceanic, land, sea-ice, and biogeochemical components on the sphere, with 25+ dynamical cores spanning five discretisation families (C-D grid FV3, finite-volume PPM, pseudospectral, FC-Gram, C-grid) and 35+ physics parameterisations. The cubed-sphere implementation is unified around a single FV3-style C-D grid discretisation (Lin 2004, Putman & Lin 2007) shared by atmosphere (shallow water, hydrostatic PE, non-hydrostatic CE) and ocean, eliminating the Hollingsworth–Kallberg instability. The cubed-sphere transport uses an FV3-faithful PPM reconstruction with monotonicity limiters and flux-form mass transport on staggered grids, complemented by an FV3-style forward–backward shallow-water core (`fv3_sw_core.py`) implementing the full `c_sw/d_sw` stepping (kernels certified bit-exact against Fortran; full-chain fidelity in progress) and `d2a2c_vect` (D-grid $\rightarrow$ A-grid $\rightarrow$ C-grid) vector conversion with duo-grid sub-grid metrics ($\sin \alpha_{sg}$, $\cos \alpha_{sg}$) following Mouallem et al. (2023). The model includes a land carbon cycle with Farquhar photosynthesis and stomatal conductance, ocean biogeochemistry (abiotic carbon + NPZD ecosystem), sea-ice dynamics with EVP rheology and multi-category ice, CMIP6-class infrastructure (table-driven CF/CMOR output across eleven CMIP6 tables, experiment templates, restart reproducibility, documented tuning), and a fully differentiable 4D-Var data assimilation system leveraging JAX autodiff for adjoint-free gradient computation. The MPI-distributed runtime supports native 4D halo exchange (one message per neighbour for all vertical levels), reverse-mode AD through MPI via a custom `sendrecv` VJP wrapper, and AD-safe `allreduce(SUM)` for gradient-compatible global reductions. The codebase comprises approximately 990 source modules ($\sim$ 590,000 lines of code, excluding the vendored CLM-ML canopy backend) exercised by roughly 32,000 test functions across some 2,700 test files.

What legoESM 1.0 adds. Version 1.0 is the first release in which a production climate configuration, rather than only the component library, is a versioned part of the model. The headline additions since the previous documentation release are:

- a **production AMIP configuration** on the MPAS Voronoi mesh with a CAM6-homogenised physics suite — CAM6 32-level hybrid table, Zhang–McFarlane convection, prognostic CLUBB with CAM6 cloud fraction, MG2-flavour two-moment microphysics, McICA cloud overlap, and an interactive multilayer land (Section 18.2);
- **conservation repairs**: on the production lane, conservative tracer vertical transport on the hybrid coordinate and column-conserving tracer positivity on every grid; in the spectral core, Simmons–Burridge vertical transport and frictional heating; and a single temperature-dependent latent-heat family used by the surface-flux interface, sea ice and the land canopy (schemes ported from a reference model, and the FV3 duo-grid lane, keep their own constants) (Sections 15.2 and 16);
- a hardened **parallel runtime** for multi-GPU and multi-node runs, with size-aware halo scheduling, space-filling-curve ordering, a communication-free ocean preconditioner, and explicit precision policies (Section 11);
- an **FV3 duo-grid lane** whose adiabatic deck is certified against the Fortran FV3 core (Section 7.5);
- table-driven **CMIP6 output** read from the official CMOR tables, with interval means and multi-rank output on the MPAS lane (Section 18);
- unified **training and calibration lanes** that train one model registry across the WeatherBench and AIMIP campaigns and refuse parameters that carry no gradient (Section 14).

Design principles.

- Pure-JAX implementation: all components are compatible with `jit`, `grad`, and `vmap`.




Figure 1: **Federation dependency DAG.** Arrows point from a package to what it imports (“depends on”); the layering is gated by zero-ignore import contracts (see **FEDERATION.md**). The four components are mutually independent and standalone-installable (contract #2); the orchestration cluster (coupler / ml / tools) sits above and carries the only mutual edges.

- Modular physics: each parameterisation follows a factory pattern returning a callable `physics_fn(state, grid, coord) → Tendencities`.
- Multiple dynamical cores share the same physics interface, enabling interchangeable experiments.
- Neural-network components (SFNO) integrate seamlessly via the same JAX computational graph.

Model Architecture

legoESM is organised as a *federation* of eight independently installable packages that all ship into one shared `legoesm` PEP-420 namespace, layered as an acyclic dependency graph whose import boundaries are enforced by zero-ignore `lint-imports` contracts (Fig. 1). A single `legoesm-core` substrate provides fields, grids, vertical coordinates, operators, the device/precision runtime, the parallel halo/reduction layer, and the grid $\times$ complexity capability matrix; the four Earth-system components (`atmosphere`, `ocean`, `land`, `ice`) depend *only* on the substrate and never on each other, so each can be `pip installed` and gradient-checked in isolation. Above them sits an orchestration cluster — `coupler+driver`, `ml+training+da`, and `tools` — whose mutual edges form a uv-workspace cycle, so the three install together.

The “lego” in legoESM is the second design axis: standard tensor-in / tendency-out interfaces let every block — dynamical core, each physics parameterisation, grid, vertical coordinate, time integrator, and even classical-versus-neural physics — be swapped independently by a string scheme name resolved through a factory (Table 1). Figure 2 traces one coupled time step through these swap points.

A coupled step runs at three nested cadences. The host *segment loop* (`build_segment_fn → run_segment_jit`) drives a `jax.lax.scan` over inner atmosphere steps; each inner step advances the dycore (u, v, T, p_s), applies the gated dry-mass fixer, runs the unified physics pipeline (radiation sub-cycled by `lax.cond`; then convection, microphysics, surface exchange, turbulence/BL,




Figure 2: Coupled time-step data flow. White blocks tagged [lego] are factory-swapped by scheme name (Table 1); dark blocks are fixed numerics. The atmosphere inner step (blue) runs inside a `jax.lax.scan`; the slow components (purple) sub-cycle at `coupling_dt`. The orange link is the bidirectional atmosphere–surface flux exchange. The entire graph executes on one `ParallelRuntime` with AD-safe halo exchange and `allreduce(SUM)` reductions, so `jax.grad` flows end-to-end through the coupled step.

Table 1: **The “lego” swap dials.** Each axis is selected independently by a string scheme name resolved through the listed factory; an unknown name raises **ValueError**. Menus abbreviated.

Swap axis	Factory entry point	Example menu
Horizontal grid	<code>create_grid()</code>	cubed-sphere, Gaussian/spectral, lat-lon, MPAS Voronoi, tripolar (ocean)
Dynamical core	<code>create_atmosphere_dycore()</code>	{cdgrid, spectral, mpas, sfno} $\times$ {shallow water, hydrostatic PE, NH compressible Euler}; lat-lon C-grid; FV3 duo-grid; plane
Vertical coordinate	<code>create_*_coordinate()</code>	sigma, hybrid σ - p , CAM6 L32 hybrid table, height / z^*
Time integrator	<code>dispatch_integrator()</code>	ssp_rk3, ssp_rk3_scan, ssp_rk34, ssp_rk54, ssp_rk54_scan, rk4
Radiation	<code>make_radiation_physics()</code>	gray (Frierson), simple LW, RRTMGP correlated- k , 3-D Monte-Carlo (plane LES/CRM only)
Convection	<code>make_convection_physics()</code>	Bechtold, Tiedtke, Kain-Fritsch, Zhang-McFarlane, Emanuel, Kuo, SBM, mass flux, EDMF, DCA
Microphysics	<code>make_microphysics_physics()</code>	Kessler, Sundqvist, Morrison (MG2 / SAM flavours), Seifert-Beheng, Thompson, P3, fast SBM, SDM, ML emulator
Turbulence / BL	<code>make_turbulence_physics()</code>	Louis, Holtslag-Boville, YSU, EDMF, TKE, MYNN-2.5, CLUBB (prognostic), CLUBB-lite, Smagorinsky
Gravity-wave drag	<code>make_gwd_physics()</code>	Lindzen, Hines, McFarlane, E3SM/CAM, Rayleigh, prognostic spectral, ML
Physics core	<code>make_{physics,hybrid,neural}</code>	classical, hybrid, neural (SFNO)
Ocean vertical mixing	<code>make_vertical_mixing_physics()</code>	KPP, Richardson, constant, TKE (NEMO <code>zdf_tke</code>), CATKE; tidal mixing applied as a separate additive step
Ocean lateral mixing	<code>make_lateral_mixing_physics()</code>	harmonic, biharmonic, GM-Redi
Ocean surface forcing	<code>make_surface_forcing_physics()</code>	bulk, restoring, prescribed
Complexity ladder	<code>resolve_model_complexity()</code>	idealized, intermediate, full

and gravity-wave drag summed as tendencies), and closes with the moisture/energy state update. After each segment, `_segment_hook` sub-cycles the slow components at `coupling_dt` (default 3600s): the ocean step, the tile surface coupler (area-weighted bulk-flux blend over ocean/ice/land/lake), and the gated CO₂ tracer, followed by conservation diagnostics. Because every block is a pure JAX pytree function and all communication is differentiable, the same graph supports forward simulation, `jit` compilation, and reverse-mode `grad`.

Grid Systems

Cubed-Sphere Grid

Six faces of a cube are mapped to the sphere via **gnomonic (central) projection** (Ronchi et al., 1996; Putman and Lin, 2007). This quasi-uniform tiling replaces the polar singularity and meridian convergence of a latitude-longitude mesh with six logically rectangular panels, so the maximum stable timestep is set by a near-uniform grid spacing rather than by the vanishing zonal spacing at the poles. Let $(\alpha_x, \alpha_y) \in [-\pi/4, \pi/4]^2$ be the equiangular gnomonic coordinates on a face; because great circles map to straight lines in (α_x, α_y) , the projection is naturally suited to conservative finite-volume discretisation (Sadourny, 1972).

Face-to-Cartesian mapping

For face 0 (+x face), the gnomonic map sends the planar face coordinates to a ray from the cube centre:

$$(x, y, z) = (1, \tan \alpha_x, \tan \alpha_y), \quad (1)$$

with analogous expressions for the remaining five faces obtained by axis permutation and sign. Points are normalised to the unit sphere: $\hat{\mathbf{r}} = \mathbf{r}/|\mathbf{r}|$, and the geographic longitude λ and latitude ϕ follow as $\lambda = \text{atan2}(\hat{y}, \hat{x})$, $\phi = \arcsin(\hat{z})$.

Cell area

The area element on the gnomonic equidistant grid is

$$\Delta A = \frac{R^2 \Delta \alpha^2}{\cos^2 \alpha_x \cos^2 \alpha_y D^3}, \quad D = \sqrt{1 + \tan^2 \alpha_x + \tan^2 \alpha_y}, \quad \Delta \alpha = \frac{\pi}{2n}. \quad (2)$$

Here R is the planetary radius, $\Delta \alpha$ the uniform angular cell width on a Cn face of $n \times n$ cells, and D the distance from the cube centre to the projected point. The D^{-3} factor is exactly the Jacobian of the gnomonic map; summing ΔA over the six faces recovers $4\pi R^2$, the discrete area-conservation property the finite-volume operators depend on.

Grid spacing

Great-circle distances use the chord-to-arc formula:

$$d = R \cdot 2 \arcsin\left(\frac{|\mathbf{r}_1 - \mathbf{r}_2|}{2}\right). \quad (3)$$

Metric factors $\Delta x_{i,j}$ and $\Delta y_{i,j}$ are computed as two-cell spans for second-order centred differences.

Grid angle and wind rotation

The angle θ between the grid x -axis and geographic east is

$$\theta = \text{atan2}\left(\frac{\partial \phi}{\partial i}, \frac{\partial \lambda}{\partial i} \cos \phi\right). \quad (4)$$

Wind rotation between geographic (u_E, v_N) and grid-aligned (u_g, v_g) :

$$\begin{pmatrix} u_E \\ v_N \end{pmatrix} = \begin{pmatrix} \cos \theta & -\sin \theta \\ \sin \theta & \cos \theta \end{pmatrix} \begin{pmatrix} u_g \\ v_g \end{pmatrix}. \quad (5)$$

Halo exchange

Scalar fields are exchanged directly across the six-face connectivity table. Vector fields are first rotated to geographic components, padded as scalars, then rotated back using the padded grid angle, because the grid-aligned basis is discontinuous across a panel edge and a naive copy of (u_g, v_g) would inject a spurious rotation at the seam. A discrete duo-grid treatment of the panel edges removes the residual edge imprint that a single-grid halo leaves on differential operators (Mouallem et al., 2023).

Two halo widths are supported. **Halo = 1** (one ghost cell per side) suffices for second-order centred differences and MUSCL reconstruction. **Halo = 2** (two ghost cells per side) is required by higher-order reconstruction stencils (PPM, WENO5). The halo exchange interpolates neighbour-face strips to the correct gnomonic position at each depth, removing the $O(\Delta \alpha)$ mismatch near cube corners.

Under MPI, the native **4D halo exchange** path (`pad_halo_4d`, `pad_halo_vector_4d`) exchanges all vertical levels in a single MPI message per neighbour, replacing the earlier per-level `vmap(pad_halo)` path. All `sendrecv` calls go through a `@jax.custom_vjp` wrapper that swaps source/dest in the backward pass, enabling `jax.grad` through distributed halo exchange for differentiable MPI simulations.

Per-step collectives are further compressed on every dycore backend: the cubed-sphere all-reduces are fused, the spectral FFT and Legendre transforms run as a single pass with stacked per-tracer analyses, and the MPAS exchanges are grouped into coloured rounds. The multi-device runtime, its halo schedules and their defaults are described in Section 11; measured scaling is documented in `docs/performance/REAL_HARDWARE_SCALING.md`.

Gaussian Grid and Spectral Transforms

The spectral dynamical cores represent fields by their spherical-harmonic coefficients and use a Gaussian grid with triangular truncation at total wavenumber $n_{\max}$ (Bourke, 1972). Triangular truncation is chosen because it is isotropic: the retained set of harmonics is invariant under rotation of the sphere, so no direction is preferentially resolved.

Grid dimensions

$$n_{\text{lat}} = \lceil \frac{3}{2}(n_{\max} + 1) \rceil_{\text{even}}, \quad n_{\text{lon}} = 2 n_{\text{lat}}, \quad n_{\text{sh}} = \frac{(n_{\max} + 1)(n_{\max} + 2)}{2}. \quad (6)$$

The factor $\frac{3}{2}$ is the standard transform-grid choice that makes the Gaussian quadrature exact for the quadratic (advective) products, eliminating spectral aliasing. Latitudes are Gaussian quadrature points $\mu_k = \sin \phi_k$ with weights w_k ; longitudes are equispaced so the zonal direction can be treated by FFT.

Associated Legendre polynomials

Fully normalised (4π) polynomials $P_n^m(\mu)$ satisfy $\int |Y_n^m|^2 dA = 1$. They are computed via the recursion

$$P_m^m = -\sqrt{\frac{2m+1}{2m}} \sqrt{1-\mu^2} P_{m-1}^{m-1}, \quad (7)$$

$$P_n^m = \sqrt{\frac{4n^2-1}{n^2-m^2}} \left(\mu P_{n-1}^m - \sqrt{\frac{(n-1)^2-m^2}{4(n-1)^2-1}} P_{n-2}^m \right). \quad (8)$$

The derivative operator $H_n^m = -(1-\mu^2) dP_n^m/d\mu$ is used for meridional gradients.

Spherical harmonic transforms

Analysis (grid $\rightarrow$ spectral).

$$\hat{f}_{nm} = 2\pi \sum_k w_k P_n^m(\mu_k) \tilde{f}_m(\mu_k), \quad \tilde{f}_m = \frac{1}{n_{\text{lon}}} \text{FFT}[f]_m. \quad (9)$$

Synthesis (spectral $\rightarrow$ grid).

$$f(\phi, \lambda) = \text{iFFT} \left[\sum_{n=m}^{n_{\max}} \hat{f}_{nm} P_n^m(\mu) \right]. \quad (10)$$

Spectral operators

Laplacian eigenvalue: $\nabla^2 \hat{f}_{nm} = -n(n+1)/a^2 \cdot \hat{f}_{nm}$. Inverse Laplacian: $\nabla^{-2} \hat{f}_{nm} = -a^2/[n(n+1)]$ for $n > 0$. Hyperdiffusion: $-\nu[n(n+1)/a^2]^p \hat{f}_{nm}$.

Velocity from vorticity and divergence

Streamfunction $\psi = \nabla^{-2} \zeta$, velocity potential $\chi = \nabla^{-2} \delta$. Velocity reconstruction:

$$u \cos \phi = \frac{\partial \psi}{\partial \theta} + \frac{\partial \chi}{\partial \lambda}, \quad v \cos \phi = \frac{\partial \psi}{\partial \lambda} - \frac{\partial \chi}{\partial \theta}, \quad (11)$$

where θ is colatitude.

Sigma Vertical Coordinate

$$\sigma = p / p_s, \quad \sigma \in [\sigma_{\text{top}}, 1]. \quad (12)$$

Half-level interfaces $\sigma_{k+1/2}$ are uniformly spaced; full levels $\sigma_k = (\sigma_{k-1/2} + \sigma_{k+1/2})/2$. Layer thickness $\Delta\sigma_k = \sigma_{k+1/2} - \sigma_{k-1/2}$.

Simmons–Burridge geopotential

The vertical finite differencing follows Simmons and Burridge (1981), constructed so that the discrete pressure-gradient and energy-conversion terms conserve total energy under adiabatic, frictionless flow. The geopotential at full level k is

$$\Phi_k = \Phi_s + R_d \sum_{k'=k+1}^{N-1} T_{k'} \ln \frac{\sigma_{k'+1/2}}{\sigma_{k'-1/2}} + \alpha_k R_d T_k, \quad (13)$$

where the Simmons–Burridge coefficient is

$$\alpha_k = 1 - \frac{\sigma_{k-1/2}}{\Delta\sigma_k} \ln \frac{\sigma_{k+1/2}}{\sigma_{k-1/2}}. \quad (14)$$

Sigma-dot

The vertical velocity in σ -space is obtained by integrating the continuity equation from the model top ($\dot{\sigma} = 0$):

$$\dot{\sigma}_{k+1/2} = \frac{\sigma_{k+1/2} - \sigma_{\text{top}}}{1 - \sigma_{\text{top}}} D_{\text{tot}} - \sum_{k'=0}^k \nabla \cdot \mathbf{v}_{k'} \Delta\sigma_{k'}, \quad D_{\text{tot}} = \sum_{k'=0}^{N-1} \nabla \cdot \mathbf{v}_{k'} \Delta\sigma_{k'}. \quad (15)$$

Pressure velocity

Since $p = \sigma p_s$, the material pressure tendency is

$$\omega \equiv \frac{dp}{dt} = \sigma \frac{dp_s}{dt} + p_s \dot{\sigma}, \quad (16)$$

where dp_s/dt is the *material* derivative of surface pressure (Eulerian tendency plus advection by the horizontal wind).

Height and z^* Coordinates

The terrain-following z^* coordinate:

$$z^* = H \frac{z - z_s}{H - z_s}, \quad J = \frac{\partial z}{\partial z^*} = \frac{H - z_s}{H}. \quad (17)$$

Physical height: $z(z^*, x, y) = z_s + z^* (H - z_s)/H$.

For the **ocean**, the z^* coordinate has a dynamic Jacobian:

$$J(\eta) = \frac{\eta + H}{H_{\max}}, \quad (18)$$

where η is the free-surface elevation and H the resting depth.

Core Operators

All finite-difference operators use second-order centred differences on the A-grid (collocated) with halo padding on the cubed sphere. The discrete difference operators δ_i and δ_j are defined as centred differences of *cell-centre* values on the padded grid:

$$\delta_i[q]_{i,j} = q_{i+1,j} - q_{i-1,j}, \quad \delta_j[q]_{i,j} = q_{i,j+1} - q_{i,j-1}. \quad (19)$$

All metric-weighted fluxes ($u h_y$, $v h_x$) are evaluated at cell centres before differencing, consistent with the A-grid stencil.

Gradient

$$\left. \frac{\partial f}{\partial x} \right|_{i,j} = \frac{f_{i+1,j} - f_{i-1,j}}{2 \Delta x_{i,j}}, \quad \left. \frac{\partial f}{\partial y} \right|_{i,j} = \frac{f_{i,j+1} - f_{i,j-1}}{2 \Delta y_{i,j}}. \quad (20)$$

Divergence (flux form)

$$\nabla \cdot (u, v) = \frac{1}{2 A_{i,j}} \left[\delta_i(u h_y) + \delta_j(v h_x) \right], \quad (21)$$

where $h_x = \Delta x/2$, $h_y = \Delta y/2$ are half-edge lengths and A is the cell area.

Vorticity

$$\zeta = \frac{1}{2 A_{i,j}} \left[\delta_i(v h_y) - \delta_j(u h_x) \right]. \quad (22)$$

C-D Grid Operators (FV3-style, `operators_cdgrid.py`)

The canonical cubed-sphere operator module uses FV3-style C-D grid staggering (Lin, 2004; Putman and Lin, 2007), shared by all atmosphere (SW, PE, CE) and ocean dynamical cores. D-grid winds at cell corners $(6, n+1, n+1)$ are prognostic; C-grid velocities at cell edges are diagnosed for mass/scalar transport. The C-D arrangement combines the favourable geostrophic adjustment of the C-grid with the cleanly defined corner vorticity of the D-grid.

FV3-faithful PPM transport (v3.9). The operator module was rewritten (~ 740 lines) with proper Piecewise Parabolic Method reconstruction (Colella and Woodward, 1984), Colella–Woodward monotonicity limiters, and flux-form mass transport on staggered grids, faithfully following the FV3 algorithm of Putman and Lin (2007). This replaces the earlier simplified A-grid transport with a fully staggered C-D grid transport that correctly handles cross-face halo exchange, geographic-to-grid wind rotation, and corner interpolation for the cubed sphere.

D-grid vorticity. Vorticity at each cell corner is the discrete circulation integral divided by the corner dual-cell area A_c :

$$\zeta_{\text{corner}} = \frac{1}{A_c} \oint \mathbf{u} \cdot d\mathbf{l} \approx \frac{1}{A_c} [u_{i,j}^D \Delta x - u_{i+1,j}^D \Delta x + v_{i,j+1}^D \Delta y - v_{i,j}^D \Delta y]. \quad (23)$$

Because vorticity is defined as the line integral of the tangential wind around the dual cell, the discrete curl inherits the Stokes-theorem identity exactly, so the kinetic-energy and enstrophy budgets remain mutually consistent. This circulation-based discretisation is exact on the D-grid and eliminates the Hollingsworth–Kallberg instability that plagues A-grid vorticity operators.

Arakawa–Lamb gradient. The Bernoulli (or Exner) gradient at D-grid corners is a 4-point Arakawa–Lamb average:

$$\left. \frac{\partial B}{\partial x} \right|_{i+\frac{1}{2},j+\frac{1}{2}} = \frac{B_{i+1,j} + B_{i+1,j+1} - B_{i,j} - B_{i,j+1}}{2\Delta x}, \quad (24)$$

and analogously for $\partial B/\partial y$.

C-grid mass-flux divergence. Scalar transport uses upwind C-grid mass fluxes:

$$\frac{dh}{dt} = -\frac{1}{A} [(hu_c)_{i+\frac{1}{2},j} - (hu_c)_{i-\frac{1}{2},j} + (hv_c)_{i,j+\frac{1}{2}} - (hv_c)_{i,j-\frac{1}{2}}], \quad (25)$$

where u_c, v_c are diagnosed from u_d, v_d by face-edge averaging.

FV3 forward–backward shallow-water core (fv3_sw_core.py). An FV3-style implementation (kernels certified bit-exact; whole-chain fidelity in progress) of the c_sw/d_sw forward–backward stepping algorithm (Lin, 2004). The forward step (c_sw) computes C-grid mass fluxes and height tendencies; the backward step (d_sw) updates D-grid winds using the vorticity flux and pressure gradient. Key components:

- **d2a2c_vect:** D-grid $\rightarrow$ A-grid $\rightarrow$ C-grid vector conversion, faithfully porting the GFDL FV3 algorithm including edge-boundary stencils, one-sided interpolation near cube edges, and non-orthogonality corrections via $\cos\alpha/\sin^{-2}\alpha$ metrics.
- **PPM wind advection:** Piecewise Parabolic Method reconstruction for C-grid wind tendencies (xtp_u, ytp_v), with FV3-specific interpolation coefficients (A_1, A_2, C_1-C_3).
- **PPM vorticity transport:** Monotonic PPM reconstruction for the absolute vorticity flux, stabilising the D-grid wind update.
- **Kinetic energy:** Cell-centred kinetic energy computed from A-grid velocities with optional 4th-order averaging via a2b_ord4.

Duo-grid sub-grid metrics (Mouallem et al., 2023). Following the duo-grid construction of Mouallem et al. (2023), the cubed-sphere C-D grid includes 9-point sub-grid angle metrics per cell, stored as $\sin\alpha_{\text{sg}}$ and $\cos\alpha_{\text{sg}}$ arrays of shape $(6, n, n, 9)$. These encode the local deviation from orthogonality at cell centres, edge midpoints, and corners, enabling FV3-faithful flux scaling at cube-face boundaries where the gnomonic projection introduces non-orthogonality. Additional C-D grid metrics include centre-to-centre distances $\Delta x_c, \Delta y_c$, and non-orthogonality factors $\cos\alpha_{\text{corner}}, (\sin^2\alpha_{\text{corner}})^{-1}$ at D-grid corners.

Laplacian and Hyperdiffusion

$$\nabla^2 f = \nabla \cdot (\nabla f), \quad \text{Hyperdiffusion:} \quad -\nu \nabla^4 f = -\nu \nabla^2 (\nabla^2 f). \quad (26)$$

Finite-Volume Transport

The finite-volume dynamical cores use upwind reconstruction at cell edges coupled with a Rusanov approximate Riemann solver. The flux-form, dimensionally split transport follows the multidimensional flux-form semi-Lagrangian construction of Lin and Rood (1996) and the vertically Lagrangian finite-volume core of Lin (2004), which together guarantee discrete mass conservation. Three reconstruction methods are available, selected via the `limiter` parameter, trading order of accuracy against the width of the halo stencil.

MUSCL Reconstruction (2nd order, halo = 1)

Piecewise-linear reconstruction with slope limiting (van Leer, 1979). The limiter clips the reconstructed slope so the edge states never overshoot the neighbouring cell averages, preserving monotonicity. At the interface between cells i and $i+1$:

$$q_L = q_i + \frac{1}{2} \phi_i, \quad q_R = q_{i+1} - \frac{1}{2} \phi_{i+1}, \quad (27)$$

where ϕ_i is the limited slope. Two limiters are provided:

- **Minmod**: $\phi = \text{minmod}(\Delta q_+, \Delta q_-)$,
- **MC** (monotonised central): $\phi = \text{minmod}(\text{minmod}(2\Delta q_+, 2\Delta q_-), \frac{1}{2}(\Delta q_+ + \Delta q_-))$.

PPM Reconstruction (4th order, halo = 2)

The Piecewise Parabolic Method (Colella and Woodward, 1984) constructs a monotone parabola in each cell from two edge values and the cell average, giving fourth-order accuracy in smooth regions while suppressing oscillations at sharp gradients.

Fourth-order edge values

At the interface $i + \frac{1}{2}$:

$$\hat{q}_{i+1/2} = \frac{7}{12}(q_i + q_{i+1}) - \frac{1}{12}(q_{i-1} + q_{i+2}). \quad (28)$$

Monotonicity constraint

The edge value is clamped to lie between its two adjacent cell averages:

$$\hat{q}_{i+1/2} \leftarrow \text{median}(\hat{q}_{i+1/2}, q_i, q_{i+1}). \quad (29)$$

Colella–Woodward limiting

For each cell i with left edge value $a_L = \hat{q}_{i-1/2}$ and right edge value $a_R = \hat{q}_{i+1/2}$:

1. If $(a_R - q_i)(q_i - a_L) \leq 0$ (local extremum): set $a_L = a_R = q_i$.
2. Define $\delta_m = a_R - a_L$ and $\delta_6 = 6(q_i - \frac{1}{2}(a_L + a_R))$.
 - If $\delta_m \delta_6 > \delta_m^2$: set $a_L = 3q_i - 2a_R$.
 - If $\delta_m \delta_6 < -\delta_m^2$: set $a_R = 3q_i - 2a_L$.

The left and right states at each edge are then $q_L = a_R$ of the cell to the left and $q_R = a_L$ of the cell to the right.

WENO5 Reconstruction (5th order, halo = 2)

The 5th-order Weighted Essentially Non-Oscillatory scheme (Jiang and Shu, 1996) combines three 3rd-order sub-stencil polynomials with nonlinear weights that adapt to local smoothness, downweighting any stencil that straddles a discontinuity so that the reconstruction stays high-order in smooth flow yet essentially non-oscillatory near shocks.

Left-biased reconstruction at $i + \frac{1}{2}$

Using cells $q_{i-2}, \dots, q_{i+2}$:

$$q^{(0)} = \frac{1}{6}(2q_{i-2} - 7q_{i-1} + 11q_i), \quad d_0 = \frac{1}{10}, \quad (30)$$

$$q^{(1)} = \frac{1}{6}(-q_{i-1} + 5q_i + 2q_{i+1}), \quad d_1 = \frac{6}{10}, \quad (31)$$

$$q^{(2)} = \frac{1}{6}(2q_i + 5q_{i+1} - q_{i+2}), \quad d_2 = \frac{3}{10}. \quad (32)$$

Smoothness indicators

$$\beta_0 = \frac{13}{12}(q_{i-2} - 2q_{i-1} + q_i)^2 + \frac{1}{4}(q_{i-2} - 4q_{i-1} + 3q_i)^2, \quad (33)$$

$$\beta_1 = \frac{13}{12}(q_{i-1} - 2q_i + q_{i+1})^2 + \frac{1}{4}(q_{i-1} - q_{i+1})^2, \quad (34)$$

$$\beta_2 = \frac{13}{12}(q_i - 2q_{i+1} + q_{i+2})^2 + \frac{1}{4}(3q_i - 4q_{i+1} + q_{i+2})^2. \quad (35)$$

Nonlinear weights

$$\alpha_k = \frac{d_k}{(\varepsilon + \beta_k)^2}, \quad \omega_k = \frac{\alpha_k}{\sum_j \alpha_j}, \quad q_L = \sum_k \omega_k q^{(k)}, \quad \varepsilon = 10^{-6}. \quad (36)$$

Right-biased reconstruction

The right state q_R at $i + \frac{1}{2}$ is obtained by mirroring the stencil around the interface, using cells $q_{i-1}, \dots, q_{i+3}$ with ideal weights $d_0 = \frac{1}{10}$, $d_1 = \frac{6}{10}$, $d_2 = \frac{3}{10}$ applied to the reversed sub-stencils.

Rusanov Flux

At each edge, the mass flux is

$$F = \frac{1}{2}(h_L u_L + h_R u_R) - \frac{1}{2} \alpha (h_R - h_L), \quad \alpha = \max(|u_L| + c_L, |u_R| + c_R), \quad (37)$$

where $c = \sqrt{g h}$ is the gravity-wave speed (set to zero for pure transport without wave-speed dissipation). The net flux divergence for cell i is

$$-(\nabla \cdot \mathbf{F})_i = -\frac{1}{A_i} \left[(F_x h_y)_{i+1/2} - (F_x h_y)_{i-1/2} + (G_y h_x)_{j+1/2} - (G_y h_x)_{j-1/2} \right], \quad (38)$$

where h_x, h_y are the edge metric factors.

Well-balanced scheme

When topography h_s is present, the scheme reconstructs the free surface $\eta = h + h_s$ instead of h . The fluid depth at each edge is recovered as $h_L = \max(\eta_L - h_{s,\text{edge}}, 0)$, ensuring zero spurious flux for a lake-at-rest state (Audusse et al., 2004).

FC-Gram Spectral Operators

The Fourier Continuation (FC-Gram) method extends periodic Fourier spectral methods to non-periodic domains by appending a smooth continuation interval. On each cubed-sphere face, the 1D function $f(x)$ on $[0, 1]$ is extended to a periodic function on $[0, 1 + d]$ via a Gram polynomial basis in the continuation region.

Extension operator

Given f sampled at n equispaced points, the extension f_{ext} is constructed as

$$f_{\text{ext}}(x) = \sum_{k=0}^{K-1} \alpha_k G_k(x), \quad x \in [1, 1+d], \quad (39)$$

where G_k are orthogonal Gram polynomials and the coefficients α_k are chosen so that f_{ext} smoothly matches f at both endpoints, yielding a globally smooth periodic extension.

Spectral differentiation

On the extended periodic domain of length $L_{\text{ext}} = 1 + d$, differentiation is performed via FFT:

$$f'_j = \text{IFFT} \left[i k \text{FFT}[f_{\text{ext}}] \right] \Big|_{x_j \in [0,1]}, \quad (40)$$

where k are the Fourier wavenumbers on the extended domain. This achieves spectral accuracy in the interior with controlled Gibbs artifacts at face boundaries.

C-grid-style divergence damping

The FC-Gram variants with divergence damping add 2nd and 4th order selective dissipation of divergent modes:

$$D_{\text{div}} = -\nu_2 \nabla(\nabla \cdot \mathbf{v}) - \nu_4 \nabla(\nabla^2(\nabla \cdot \mathbf{v})), \quad (41)$$

applied to the momentum equations. This selectively damps gravity-wave modes while preserving rotational (balanced) flow, mirroring the divergence damping strategy used in FV3 on its C-D grid.

Atmospheric Dynamics

Shallow Water Equations

The shallow water system is the canonical proving ground for horizontal dynamical cores: it retains the rotational and gravity-wave dynamics of the full equations while collapsing the vertical structure to one layer, so the standard test set of Williamson et al. (1992) can isolate grid-imprinting, dispersion, and conservation errors. In **vector-invariant form** on the cubed sphere:

$$\frac{\partial h}{\partial t} = -\nabla \cdot (h \mathbf{v}), \quad (42)$$

$$\frac{\partial u}{\partial t} = (\zeta + f) v - \frac{\partial B}{\partial x} + D_u, \quad (43)$$

$$\frac{\partial v}{\partial t} = -(\zeta + f) u - \frac{\partial B}{\partial y} + D_v, \quad (44)$$

where h is the fluid depth, $\mathbf{v} = (u, v)$ the horizontal velocity, f the Coriolis parameter, h_s the bottom topography, the Bernoulli function is $B = \frac{1}{2}(u^2 + v^2) + g(h + h_s)$, ζ is the relative vorticity, and D_u, D_v represent hyperdiffusion. Equation (42) expresses exact mass (volume) conservation in flux form; the momentum equations express the balance between the absolute-vorticity ($\zeta + f$) flux, the Bernoulli gradient that combines kinetic energy and free-surface pressure, and scale-selective dissipation.

The mass equation (42) is discretised using the finite-volume Rusanov flux (Section 5) with a configurable reconstruction method: MUSCL (2nd order), PPM (4th order), or WENO5 (5th order), selected via `ShallowWaterConfig.fv_limiter`. Momentum uses centred differences with velocity-only biharmonic hyperdiffusion.

Hydrostatic Primitive Equations (σ -coordinates)

Prognostic variables: u, v (winds), T (temperature), p_s (surface pressure). The vertical discretisation uses the energy- and angular-momentum-conserving finite differences of Simmons and Burridge (1981).

$$\frac{\partial p_s}{\partial t} = - \int_0^1 \nabla \cdot (p_s \mathbf{v}) d\sigma, \quad (45)$$

$$\frac{\partial u}{\partial t} = (\zeta + f) v - \frac{\partial B}{\partial x} - R_d T \frac{\partial \ln p_s}{\partial x} + D_u, \quad (46)$$

$$\frac{\partial v}{\partial t} = -(\zeta + f) u - \frac{\partial B}{\partial y} - R_d T \frac{\partial \ln p_s}{\partial y} + D_v, \quad (47)$$

$$\frac{\partial T}{\partial t} = -\nabla \cdot (T \mathbf{v}) + T \nabla \cdot \mathbf{v} - \dot{\sigma} \frac{\partial T}{\partial \sigma} + \kappa T \frac{\omega}{p} + Q, \quad (48)$$

where $B = K + \Phi$, $K = \frac{1}{2}(u^2 + v^2)$, Φ is the geopotential, $\dot{\sigma}$ the vertical velocity in σ -space, $\omega = dp/dt$ the pressure velocity, Q the diabatic heating from physics, and $\kappa = R_d/c_p$. The surface-pressure tendency (45) is the vertically integrated mass continuity that closes the hydrostatic system: with vertical momentum replaced by hydrostatic balance, it is column mass divergence—not vertical acceleration—that sets p_s . The $R_d T \nabla \ln p_s$ term in the momentum equations accounts for the tilt of the σ surfaces, and the adiabatic term $\kappa T \omega/p$ in (48) represents compression warming during descent.

Spectral Primitive Equations (Vorticity–Divergence)

On the Gaussian grid, the momentum equations are recast in the vorticity–divergence form (Bourke, 1972), in which the prognostic scalars $\hat{\zeta}$ and $\hat{\delta}$ replace the vector momentum and the horizontal Laplacian is diagonal in spectral space:

$$\frac{\partial \hat{\zeta}}{\partial t} = \nabla \times [\widehat{[(\zeta + f)\mathbf{v}]}] + \widehat{\text{vert. adv.}} - \nu \nabla^{2p} \hat{\zeta}, \quad (49)$$

$$\frac{\partial \hat{\delta}}{\partial t} = \nabla \cdot [\widehat{[(\zeta + f)\mathbf{v}]}] - \nabla^2 (K + \Phi + R_d T \ln p_s) + \widehat{\text{vert. adv.}} - \nu \nabla^{2p} \hat{\delta}, \quad (50)$$

$$\frac{\partial \hat{T}}{\partial t} = -\nabla \cdot (\widehat{T\mathbf{v}}) + T \nabla \cdot \mathbf{v} - \dot{\sigma} \widehat{\partial T / \partial \sigma} + \kappa \widehat{T\omega/p}, \quad (51)$$

$$\frac{\partial \widehat{\ln p_s}}{\partial t} = - \int_0^1 \hat{\delta} d\sigma. \quad (52)$$

The spectral curl and divergence of flux vectors $(U, V) = (F_u \cos \phi, F_v \cos \phi)$ are

$$\widehat{\text{curl}} = \frac{im}{a} \mathcal{S}[V/\cos^2 \phi] + \frac{1}{a} \mathcal{S}_\mu[U], \quad \widehat{\text{div}} = \frac{im}{a} \mathcal{S}[U/\cos^2 \phi] - \frac{1}{a} \mathcal{S}_\mu[V], \quad (53)$$

where $\mathcal{S}$ and $\mathcal{S}_\mu$ denote the SH analysis with standard and $dP_n^m/d\mu$ Legendre matrices.

Energy-conserving vertical transport and frictional heating. Vertical advection of u, v and T uses the centred Simmons and Burridge (1981) pairing on both the sigma and the hybrid coordinate (`vertical_advection_scheme="sb_centered"`, the default), so that the discrete vertical transport is consistent with continuity and, in the adiabatic frictionless limit, conserves total energy. The kinetic energy removed by the explicit vorticity/divergence hyperdiffusion is returned to the temperature equation as local frictional heating (`frictional_heating=True`, default), computed from the dealiased tendencies. The legacy first-order upwind form remains selectable for reproducing older results but emits a warning, because it removes heat from the column with no physical counterpart. Tracers keep upwind vertical advection on this core: a centred scheme would overshoot and drive moisture negative.

Non-Hydrostatic Compressible Euler Equations

In the z^* terrain-following coordinate, the prognostic variables are u, v, w, θ', ρ' (perturbations from a reference state ρ_0, θ_0, π_0). Retaining the full vertical momentum equation (56) admits acoustic and non-hydrostatic gravity-wave modes, which are advanced on a fast time level by the split-explicit integrator below (Skamarock and Klemp, 2008).

$$\frac{\partial u}{\partial t} = (\zeta + f) v - \frac{\partial K}{\partial x} - c_p \theta \frac{\partial \pi'}{\partial x} + D_u - r u, \quad (54)$$

$$\frac{\partial v}{\partial t} = -(\zeta + f) u - \frac{\partial K}{\partial y} - c_p \theta \frac{\partial \pi'}{\partial y} + D_v - r v, \quad (55)$$

$$\frac{\partial w}{\partial t} = -\frac{c_p \theta}{J} \frac{\partial \pi'}{\partial z^*} - g \frac{\theta'}{\theta_0} - r w, \quad (56)$$

$$\frac{\partial \theta'}{\partial t} = -\nabla \cdot (\theta \mathbf{v}_h) - \frac{w}{J} \frac{\partial \theta}{\partial z^*} + D_\theta, \quad (57)$$

$$\frac{\partial \rho'}{\partial t} = -\frac{1}{J} \left[\nabla_h \cdot (J \rho \mathbf{v}_h) + \frac{\partial(\rho w)}{\partial z^*} \right], \quad (58)$$

where $r(z)$ is the sponge-layer Rayleigh damping coefficient and the Exner function perturbation is

$$\pi' = \pi_0 \left[\left(1 + \frac{\rho'}{\rho_0} \right) \left(1 + \frac{\theta'}{\theta_0} \right) \right]^{R_d/c_v} - \pi_0. \quad (59)$$

Split-explicit time integration

The equations are integrated with an outer SSP-RK3 scheme for the slow tendencies (advection, Coriolis, horizontal pressure gradient, diffusion) and inner forward-backward acoustic substeps for the fast variables (w, θ', ρ').

Forward step (w).

$$w^{n+1} = w^n + \Delta t_s \left[-\frac{c_p \theta}{J} \frac{\partial \pi'}{\partial z^*} - g \frac{\theta'}{\theta_0} \right]. \quad (60)$$

Backward step (ρ', θ').

$$\rho'^{n+1} = \rho'^n - \Delta t_s \frac{\partial(\rho w^{n+1})}{\partial z^*}, \quad \theta'^{n+1} = \theta'^n - \Delta t_s \frac{w^{n+1}}{J} \frac{\partial \theta}{\partial z^*}. \quad (61)$$

An optional **semi-implicit acoustic** solve replaces the forward Euler for w with a tridiagonal system to remove the vertical acoustic CFL constraint.

Acoustic off-centering (v3.7). To damp vertically propagating acoustic/gravity wave noise in long integrations, the density update can be off-centred (Skamarock and Klemp, 2008):

$$\rho_\beta'^{n+1} = (1 + \beta) \rho'^{n+1} - \beta \rho'^n, \quad (62)$$

where $\beta \geq 0$ is the off-centering parameter (`acoustic_off_centering` in the config). For $\beta = 0$ the scheme is centred (neutral); $\beta = 0.1$ is recommended for AMIP/CMIP runs. This adds a small amount of temporal diffusion to the fast modes without tightening the horizontal CFL constraint.

Sponge layer

$$r(z) = r_{\max} \sin^2 \left(\frac{\pi}{2} \frac{\max(0, z - z_{\text{bot}})}{z_{\text{top}} - z_{\text{bot}}} \right). \quad (63)$$

Turbulence closure: CRM, LES, and DNS

The diffusion terms D_u, D_v, D_w, D_θ are evaluated in flux form, $D_\phi = \rho_0^{-1} \partial_j (\rho_0 K_\phi \partial_j \phi)$, over the horizontal directions and (optionally) a mass-weighted vertical leg. The same discrete operator spans three turbulence-resolving regimes that differ only in the viscosity K_m and the grid spacing Δx (heat and scalars use $K_h = K_m/\text{Pr}$), selected by the `turbulence_closure` parameter:

- **CRM** (convection-permitting, $\Delta x \sim 1\text{--}4\text{ km}$; `smagorinsky`): the 3-D Smagorinsky–Lilly eddy viscosity (Smagorinsky, 1963) $K_m = (C_s \Delta)^2 \sqrt{\max(0, |S|^2 - \text{Pr} N^2)}$ with $\Delta = (\Delta x \Delta y \Delta z)^{1/3}$, the full strain $|S|^2 = 2S_{ij}S_{ij}$, $C_s \approx 0.19$, and the SAM `dosmagor` moist stratification cutoff N^2 . The energy-containing eddies are entirely sub-grid; this is the regime validated against gSAM for the GATE, LBA, and RCE cases.
- **LES** (large-eddy simulation, $\Delta x \sim 10\text{--}100\text{ m}$; `smagorinsky`): the same closure with $C_s \approx 0.15$ and an optional near-wall limiter $\ell_m = \min(C_s \Delta, \kappa z)$; the energy-containing eddies are resolved. gSAM’s GATE and LBA references are themselves 100 m LES, so this mode uses the identical closure at the identical resolution.
- **DNS** (direct numerical simulation, $\Delta x \rightarrow$ the Kolmogorov scale; `molecular`): no turbulence model; K_m is the constant molecular kinematic viscosity ν with $K_h = \nu/\text{Pr}$ ($\text{Pr} \approx 0.71$ for air), giving the molecular dissipation $\nu \nabla^2 \mathbf{u}$ directly.

Because the three regimes share one discrete flux-form operator, LES and DNS are configuration changes rather than separate code paths; the molecular closure reproduces the analytical viscous decay $\partial_t \hat{u} = -\nu k^2 \hat{u}$ of an isolated shear mode to within the time-integration error.

FV3 Duo-Grid Lane

Alongside the native C-D grid cores, legoESM carries a JAX twin of the FV3 `fv_dynamics` step on the six-face duo-grid cube (Mouallem et al., 2023) (`dycore.discretization="fv3_duo"`). Its kernels, halo exchange and multi-step trajectories are scored against a pinned build of the Fortran FV3 core on both the hydrostatic and non-hydrostatic arms, and FV3’s own physical constants are used on this lane so that the comparison is not confounded by the model-wide constant set. The certified deck is adiabatic with humidity as a passenger; a moist arm passes FV3’s virtual-temperature factor into the step and is selected automatically when Kessler microphysics is attached. The lane runs in float64 storage by default; reduced-precision storage is declared but refused at construction until every workspace follows the storage dtype.

The step is face-sharded across devices with ring exchanges of width proportional to the halo, and a window layout splits each face into $k_t \times k_t$ windows ($6k_t^2$ devices, `dycore.fv3_duo_windows`) with a measured per-deck halo pad that must be given explicitly. A column-model bridge (`dycore.fv3_duo_column_lane`) presents the duo state as cell columns so that the MPAS-lane physics suite can run on FV3 dynamics, applying physics increments in the FV3 manner (winds through the D-grid lift, water tracers and layer mass through the moist update) without an additive surface-pressure fixer.

Atmospheric Physics

Gray Radiation (Frierson 2006)

The gray-radiation scheme (Frierson et al., 2006) replaces the spectral detail of the longwave with a single broadband optical depth that is a prescribed function of pressure and latitude plus a water-vapour contribution, giving an inexpensive but physically self-consistent radiative driving for aquaplanet and idealized-climate experiments.

Longwave optical depth

$$\tau(\sigma, \phi) = [f_l \sigma + (1 - f_l) \sigma^4] [\tau_e + (\tau_p - \tau_e) \sin^2 \phi] + \sum_k \tau_{\text{moist}} q_{v,k} \frac{\Delta p_k}{g}, \quad (64)$$

where the moisture optical depth per layer is $\Delta \tau_{\text{moist},k} = \tau_{\text{moist}} q_{v,k} \Delta p_k / g$, with $\tau_e = 7.2$, $\tau_p = 1.8$, $f_l = 0.2$, $\tau_{\text{moist}} = 0.0115 \text{ m}^2 \text{ kg}^{-1}$.

Two-stream LW fluxes

Layer transmittance and emissivity:

$$\mathcal{T}_k = e^{-D \Delta \tau_k}, \quad \varepsilon_k = 1 - \mathcal{T}_k, \quad D = 1.66. \quad (65)$$

Upward sweep from surface ($F_{\text{sfc}}^\uparrow = \varepsilon_s \sigma_{\text{SB}} T_s^4$):

$$F_k^\uparrow = \mathcal{T}_k F_{k+1}^\uparrow + \varepsilon_k \sigma_{\text{SB}} T_k^4. \quad (66)$$

Downward sweep from TOA ($F_0^\downarrow = 0$):

$$F_k^\downarrow = \mathcal{T}_k F_{k-1}^\downarrow + \varepsilon_k \sigma_{\text{SB}} T_k^4. \quad (67)$$

Shortwave (Beer–Lambert)

$$F_{\text{SW}}^\downarrow(k) = \frac{S_0}{4} e^{-\tau_{\text{SW},0} \sigma_k}, \quad F_{\text{SW}}^\uparrow(k) = \alpha_s F_{\text{SW}}^\downarrow(\text{sfc}), \quad (68)$$

with $\tau_{\text{SW},0} = 0.22$, $S_0 = 1360 \text{ W m}^{-2}$, $\alpha_s = 0.31$. This follows the Frierson/Isca convention used in the code: only the downward SW beam is absorbed by the atmosphere, while reflected upward SW escapes directly to space.

Heating rate

Defining $F_{\text{net}}^\uparrow = F^\uparrow - F^\downarrow$ (upward positive) at each interface, the heating rate for layer k is

$$\left. \frac{dT}{dt} \right|_{\text{rad},k} = \frac{g}{c_p} \frac{F_{\text{net},k+1/2}^\uparrow - F_{\text{net},k-1/2}^\uparrow}{p_{k+1/2} - p_{k-1/2}}, \quad (69)$$

where $k + \frac{1}{2}$ is the bottom interface and $k - \frac{1}{2}$ the top (pressure increasing downward). Positive divergence of upward flux warms the layer.

RRTMGP Correlated- k Radiation

The RRTMGP backend (Mlawer et al., 1997; Pincus et al., 2019) calls JAX-RRTMGP with correlated- k gaseous optics (128 LW and 112 SW g -points), a two-stream RTE solver, and optional cloud optics. The correlated- k method reorders the line-by-line absorption spectrum by opacity so a handful of quadrature points reproduce the broadband fluxes at a fraction of the cost. Default trace-gas concentrations: $\text{CO}_2=415 \text{ ppmv}$, $\text{CH}_4=1900 \text{ ppbv}$, $\text{N}_2\text{O}=332 \text{ ppbv}$.

SBM Convection (Frierson 2007)

The simplified Betts–Miller (SBM) scheme (Frierson, 2007) relaxes the column toward a moist-adiabatic reference profile whenever it is convectively unstable, conserving column moist enthalpy. Symbols below: Γ_m is the moist-adiabatic lapse rate, q_s the saturation mixing ratio, CAPE the convective available potential energy, $\mathcal{F}$ a smooth trigger, and τ_c the relaxation timescale.

Moist adiabatic lapse rate

$$\Gamma_m = \frac{R_d T}{c_p p} \frac{1 + L_v q_s / (R_d T)}{1 + L_v^2 q_s / (c_p R_v T^2)}, \quad (70)$$

integrated upward via a trapezoidal predictor–corrector.

Saturation mixing ratio (Tetens)

$$e_s = 611.2 \exp\left(\frac{17.67 T_c}{T_c + 243.5}\right), \quad q_s = \frac{\varepsilon e_s}{p - e_s}, \quad \varepsilon = R_d / R_v \approx 0.622. \quad (71)$$

CAPE and trigger

$$\text{CAPE} = R_d \sum_k \max(0, T_{\text{parcel},k} - T_{\text{env},k}) \frac{\Delta p_k}{p_k}. \quad (72)$$

A smooth sigmoid trigger activates convection:

$$\mathcal{F} = \sigma(\kappa_s (\text{CAPE} - \text{CAPE}_c)), \quad \kappa_s = 0.01 \text{ (J/kg)}^{-1}. \quad (73)$$

Enthalpy-conserving relaxation

The reference profile T_{ref} is obtained by adjusting the moist adiabat so that column-integrated moist enthalpy is conserved:

$$\sum_k [c_p (T_{\text{ref},k} - T_k) + L_v (q_{\text{ref},k} - q_{v,k})] \Delta p_k = 0. \quad (74)$$

Newton iteration (2 steps) solves for the temperature offset. Tendencies:

$$\frac{dT}{dt} = \mathcal{F} \frac{T_{\text{ref}} - T}{\tau_c}, \quad \frac{dq_v}{dt} = \mathcal{F} \frac{q_{\text{ref}} - q_v}{\tau_c}, \quad \tau_c = 7200 \text{ s}. \quad (75)$$

Precipitation

$$P = -\frac{1}{g} \sum_k \frac{dq_v}{dt} \Delta p_k. \quad (76)$$

Held–Suarez Forcing

The Held and Suarez (1994) benchmark replaces the full physics with Newtonian relaxation of temperature toward a prescribed radiative equilibrium and Rayleigh friction in the boundary layer, providing a reproducible dynamical-core intercomparison that bypasses parameterisation uncertainty.

Equilibrium temperature

$$T_{\text{eq}} = \max\left(200, [315 - \Delta T_y \sin^2 \phi - \Delta \theta_z \ln(p/p_0) \cos^2 \phi] (p/p_0)^\kappa\right), \quad (77)$$

with $\Delta T_y = 60 \text{ K}$, $\Delta \theta_z = 10 \text{ K}$.

Newtonian relaxation

$$\frac{dT}{dt} = -k_T(\sigma, \phi) (T - T_{\text{eq}}), \quad k_T = k_a + (k_s - k_a) \max\left(0, \frac{\sigma - \sigma_b}{1 - \sigma_b}\right) \cos^4 \phi, \quad (78)$$

with $k_a = 1/(40 \text{ day})$, $k_s = 1/(4 \text{ day})$, $\sigma_b = 0.7$.

Rayleigh friction

$$\frac{d\mathbf{v}}{dt} = -k_v(\sigma) \mathbf{v}, \quad k_v = k_f \max\left(0, \frac{\sigma - \sigma_b}{1 - \sigma_b}\right), \quad k_f = 1/(1 \text{ day}). \quad (79)$$

Louis (1979) Turbulence

The Louis (1979) scheme parameterises vertical eddy fluxes through a mixing length and Richardson-number-dependent stability functions, giving a smooth, differentiable closure for the boundary-layer diffusivities K_m and K_h .

Mixing length

$$l = \frac{\kappa_{vK} z}{1 + \kappa_{vK} z/l_{\max}}, \quad \kappa_{vK} = 0.4, \quad l_{\max} = 100 \text{ m}. \quad (80)$$

Richardson number

$$\text{Ri} = \frac{N^2}{S^2}, \quad N^2 = \frac{g}{\theta_v} \frac{\partial \theta_v}{\partial z}, \quad S^2 = \left(\frac{\partial u}{\partial z}\right)^2 + \left(\frac{\partial v}{\partial z}\right)^2. \quad (81)$$

Stability functions

$$f_m(\text{Ri}) = \begin{cases} 1 - \frac{2b \text{ Ri}}{1 + 3bc l^2 |\text{Ri}|^{1/2} / \Delta z^2} & \text{Ri} < 0, \\ \frac{1}{1 + 2b \text{ Ri} / \sqrt{1 + d \text{ Ri}}} & \text{Ri} \geq 0, \end{cases} \quad (82)$$

with $b = c = d = 5$. Eddy diffusivities: $K_m = l^2 S f_m$, $K_h = K_m$.

Vertical Diffusion (Implicit Thomas Algorithm)

The backward-Euler system $-a_k \varphi_{k-1} + b_k \varphi_k - c_k \varphi_{k+1} = d_k$ is solved with the Thomas algorithm.

Tridiagonal coefficients.

$$a_k = \frac{\Delta t K_{k-1/2}}{\rho_k \Delta z_k \Delta z_{k-1/2}}, \quad c_k = \frac{\Delta t K_{k+1/2}}{\rho_k \Delta z_k \Delta z_{k+1/2}}, \quad b_k = 1 + a_k + c_k. \quad (83)$$

Forward sweep.

$$w_k = \frac{a_k}{b'_{k-1}}, \quad b'_k = b_k - w_k c_{k-1}, \quad d'_k = d_k + w_k d'_{k-1}. \quad (84)$$

Back substitution.

$$\varphi_k = \frac{d'_k + c_k \varphi_{k+1}}{b'_k}. \quad (85)$$

Surface Fluxes

Three bulk flux schemes are available, selected per tile.

Constant coefficients (default)

$$\tau_x = -\rho C_D |\mathbf{V}| u, \quad \tau_y = -\rho C_D |\mathbf{V}| v, \quad (86)$$

$$\text{SH} = \rho c_p C_H |\mathbf{V}| (T_s - T_a), \quad (87)$$

$$\text{LH} = \rho L_v C_E |\mathbf{V}| (q_s - q_a), \quad (88)$$

with $C_D = C_H = C_E = 1.5 \times 10^{-3}$ and $|\mathbf{V}| = \sqrt{u^2 + v^2 + 10^{-4}}$.

Monin–Obukhov Similarity Theory (MOST)

Stability-dependent transfer coefficients via iterative Obukhov-length computation using the Businger–Dyer flux–profile relationships (Businger et al., 1971).

Obukhov length.

$$L = -\frac{u_*^3 T_v}{\kappa_{\text{vK}} g (\theta_* + 0.61 T_a q_*)}, \quad (89)$$

where u_* , θ_* , q_* are friction velocity, temperature scale, and moisture scale.

Stability parameter. $\zeta = z_{\text{ref}}/L$.

Businger–Dyer stability functions.

$$\psi_m(\zeta) = \begin{cases} 2 \ln\left(\frac{1+x}{2}\right) + \ln\left(\frac{1+x^2}{2}\right) - 2 \arctan x + \frac{\pi}{2}, & \zeta < 0, \quad x = (1 - 16\zeta)^{1/4}, \\ -5\zeta, & \zeta \geq 0, \end{cases} \quad (90)$$

$$\psi_h(\zeta) = \begin{cases} 2 \ln\left(\frac{1+y}{2}\right), & \zeta < 0, \quad y = (1 - 16\zeta)^{1/2}, \\ -5\zeta, & \zeta \geq 0. \end{cases} \quad (91)$$

Transfer coefficients.

$$C_D = \left[\frac{\kappa_{\text{vK}}}{\ln(z/z_0) - \psi_m(\zeta)} \right]^2, \quad C_H = \frac{\kappa_{\text{vK}}^2}{[\ln(z/z_0) - \psi_m(\zeta)][\ln(z/z_{0t}) - \psi_h(\zeta)]}. \quad (92)$$

Iterative solver. The system $u_*, \theta_*, q_*, z_0, z_{0t}, z_{0q}, L$ is solved iteratively (default 5 iterations) using `jax.lax.fori_loop` for full JAX differentiability.

COARE 3.0 roughness (Fairall et al. 2003)

Over ocean, the COARE 3.0 bulk algorithm (Fairall et al., 2003) updates the roughness lengths each iteration from the evolving friction velocity:

$$z_0 = \alpha_c \frac{u_*^2}{g} + 0.11 \frac{\nu}{u_*}, \quad z_{0t} = \min(1.15 \times 10^{-4}, 5.5 \times 10^{-5} \text{Re}_*^{-0.6}), \quad (93)$$

where $\alpha_c = 0.011$ is the Charnock parameter and $\text{Re}_* = u_* z_0 / \nu$.

Large & Yeager (2004) roughness

The Large and Yeager (2004) bulk formulae specify an empirical neutral drag coefficient as a function of the neutral 10-m wind speed U_{10N} :

$$C_{DN} \times 10^3 = \frac{2.7}{U_{10N}} + 0.142 + 0.0764 U_{10N}, \quad (94)$$

with roughness recovered as $z_0 = 10 / \exp(\kappa_{\text{vK}} / \sqrt{C_{DN}})$. Stanton and Dalton numbers follow the stability-corrected LY04 parameterisation.

Additional Turbulence Schemes

Smagorinsky

Constant eddy diffusivity baseline: $K_m = K_h = \text{const.}$

YSU (Yonsei University)

Nonlocal K-profile boundary-layer scheme with explicit entrainment treatment at the PBL top.

CLUBB-lite

Higher-order prognostic moment closure with five prognostic variables ($\overline{w'^2}$, $\overline{\theta_l'^2}$, $\overline{r_t'^2}$, $\overline{w'\theta_l'}$, $\overline{w'r_t'}$) and cloud fraction diagnosis via assumed bivariate Gaussian PDF.

EDMF Turbulence

Eddy-diffusivity mass-flux unified PBL framework combining diffusive and convective transport in a single scheme.

Additional Convection Schemes

Kuo (1974)

Column moisture-excess scheme (Kuo, 1974) partitioning excess moisture between heating (fraction α_{heat}) and moistening ($1 - \alpha_{\text{heat}}$).

Mass-Flux Convection (Arakawa–Wu Type)

Entraining-detraining plume model with bulk mass flux, CAPE closure, and configurable entrainment/detrainment rates.

EDMF Convection

Combined eddy-diffusivity and mass-flux framework for shallow and deep convection.

Deep Convective Adjustment (DCA)

Bottom-to-top adjustment with smooth sigmoid triggers for JAX differentiability.

Gravity Wave Drag

Six GWD schemes are available:

- **Rayleigh friction:** Simple linear damping above a threshold height.
- **Lindzen (1981):** Orographic GWD with critical-level absorption (Lindzen, 1981).
- **McFarlane (1987):** Enhanced orographic scheme with improved saturation (McFarlane, 1987).
- **Hines (1997):** Doppler-spread parameterisation for non-orographic GWD (Hines, 1997).
- **Prognostic Spectral:** Multi-azimuthal, multi-wavenumber spectrum with prognostic wave flux and vertical propagation via `jax.lax.scan`.
- **ML Emulator:** Equinox MLP surrogate for learned GWD.

Cloud Microphysics

Five microphysics schemes:

- **Kessler (1969)**: Warm-rain autoconversion with threshold cloud water (Kessler, 1969).
- **Seifert–Beheng (2001)**: Two-moment warm rain with prognostic number concentration (Seifert and Beheng, 2001).
- **Morrison (2005)**: Double-moment ice + liquid scheme (Morrison et al., 2005).
- **Thompson**: Hybrid-moment scheme with ice crystal habits (Thompson et al., 2008).
- **Sundqvist (1988)** and **Xu–Randall (1996)** cloud fraction diagnostics (Sundqvist et al., 1989; Xu and Randall, 1996).

`cloud_fraction.compute_cloud_properties` dispatches on `config.scheme` and accepts only "sundqvist" or "xu_randall"; any other literal raises `ValueError` (no silent fallback). To disable cloud diagnostics, set `cloud_scheme="none"` on the caller's config so that `compute_cloud_properties` is skipped entirely upstream; "none" is not itself a valid scheme inside the dispatch.

Ocean Dynamics

Boussinesq Hydrostatic Primitive Equations

The ocean solves the hydrostatic Boussinesq primitive equations in the finite-volume, split-explicit free-surface formulation of Marshall et al. (1997), in which density variations are retained only in the buoyancy term and the fast external (barotropic) mode is integrated on a short time level. In vector-invariant form on the cubed sphere:

$$\frac{\partial u}{\partial t} = (\zeta + f)v - \frac{\partial K}{\partial x} - \frac{1}{\rho_0} \frac{\partial p'}{\partial x} + A_h \nabla^2 u + \frac{\partial}{\partial z} \left(A_v \frac{\partial u}{\partial z} \right) - w \frac{\partial u}{\partial z}, \quad (95)$$

$$\frac{\partial v}{\partial t} = -(\zeta + f)u - \frac{\partial K}{\partial y} - \frac{1}{\rho_0} \frac{\partial p'}{\partial y} + A_h \nabla^2 v + \frac{\partial}{\partial z} \left(A_v \frac{\partial v}{\partial z} \right) - w \frac{\partial v}{\partial z}. \quad (96)$$

The Coriolis force is split: depth-mean $f\mathbf{v}$ in the barotropic substeps, baroclinic shear $f\mathbf{v}'$ in the slow tendencies.

Tracer equations (skew-symmetric)

$$\frac{\partial \phi}{\partial t} = \frac{1}{2} \left[-\mathbf{v} \cdot \nabla \phi - \nabla \cdot (\phi \mathbf{v}) \right] - w \frac{\partial \phi}{\partial z} + K_h \nabla^2 \phi + \frac{\partial}{\partial z} \left(K_v \frac{\partial \phi}{\partial z} \right). \quad (97)$$

Free-surface equation

$$\frac{\partial \eta}{\partial t} = - \sum_k \nabla \cdot (h_k \mathbf{v}_k). \quad (98)$$

Barotropic Solver (Forward–Backward, Semi-Implicit Coriolis)

The barotropic system is solved with split-explicit forward–backward substeps:

$$\frac{\partial \eta}{\partial t} = -\nabla \cdot (H \bar{\mathbf{V}}), \quad \frac{\partial \bar{\mathbf{V}}}{\partial t} = f \hat{\mathbf{k}} \times \bar{\mathbf{V}} - g \nabla \eta + \mathbf{F}_{\text{slow}}. \quad (99)$$

Semi-implicit Coriolis. With $\alpha = \frac{1}{2} f \Delta t_s$:

$$\bar{U}^{n+1} = \frac{R_u + \alpha R_v}{1 + \alpha^2}, \quad \bar{V}^{n+1} = \frac{R_v - \alpha R_u}{1 + \alpha^2}, \quad (100)$$

where $R_u = \bar{U}^n + \alpha \bar{V}^n - \Delta t_s g \partial \eta / \partial x$ and $R_v = \bar{V}^n - \alpha \bar{U}^n - \Delta t_s g \partial \eta / \partial y$.

Equation of State (Wright 1997)

The ocean prognostic temperature variable T is **potential temperature** [degC]. The Wright (1997) equation of state computes in-situ density directly from potential temperature θ , salinity S [PSU], and pressure p [Pa]:

$$\rho = \frac{p + p_0(\theta, S)}{\lambda(\theta, S) + \alpha_0(\theta, S)(p + p_0)}, \quad (101)$$

where

$$\alpha_0 = a_0 + a_1\theta + a_2S, \quad (102)$$

$$p_0 = (b_0 + b_4S) + \theta(b_1 + \theta(b_2 + b_3\theta) + b_5S), \quad (103)$$

$$\lambda = (c_0 + c_4S) + \theta(c_1 + \theta(c_2 + c_3\theta) + c_5S). \quad (104)$$

Coefficients follow the MOM6 implementation. The density perturbation is $\rho' = \rho(\theta, S, p) - \rho_0$.

Linear Equation of State

A configurable linear EOS is also available for idealised experiments:

$$\rho = \rho_0 [1 - \alpha_T(T - T_{\text{ref}}) + \beta_S(S - S_{\text{ref}})], \quad (105)$$

with typical values $\alpha_T = 2 \times 10^{-4} \text{ K}^{-1}$, $\beta_S = 7.4 \times 10^{-4} \text{ PSU}^{-1}$. The linear EOS is pressure-independent (no hydrostatic pressure coupling iterations required), enabling faster compilation and execution. It is selected via `eos="linear"` in the ocean configuration, with parameters set via a `LinearEOSConfig` NamedTuple. A dispatcher `make_eos_fn()` returns the appropriate EOS callable.

Hydrostatic Pressure and Vertical Velocity

Baroclinic pressure by downward integration:

$$p'_k = \rho_0 g \eta + \sum_{k'=0}^{k-1} \rho'_{k'} g \Delta z_{k'} + \frac{1}{2} \rho'_k g \Delta z_k. \quad (106)$$

Vertical velocity diagnosed bottom-up from $w(\text{floor}) = 0$:

$$w_{k-1/2} = - \sum_{k'=k}^{N-1} \nabla \cdot (h_{k'} \mathbf{v}_{k'}). \quad (107)$$

Single-level barotropic mode ($N_{\text{lev}} = 1$). The z-star coordinate supports $N_{\text{lev}} = 1$, reducing exactly to the nonlinear shallow-water equations: $h = H + \eta$, with no baroclinic modes. Vertical mixing operators return zero for $N_{\text{lev}} < 2$. This enables MITgcm-style barotropic gyre experiments with a single depth-integrated layer.

Ocean Vertical Mixing

Three vertical mixing schemes are available:

K-Profile Parameterisation (KPP, Large et al. 1994)

The K-profile parameterisation (Large et al., 1994) diagnoses the boundary-layer depth h_{bl} from the bulk Richardson number criterion $\text{Ri}_b(h) = \text{Ri}_c = 0.3$. Within the boundary layer, diffusivity follows a cubic shape function:

$$K(z) = h_{\text{bl}} w_* G(\sigma), \quad G(\sigma) = \sigma(1 - \sigma)^2, \quad \sigma = -z/h_{\text{bl}}, \quad (108)$$

where w_* is a turbulent velocity scale incorporating wind and buoyancy forcing. Non-local transport of temperature and salinity: $\gamma_T = -C_s (\overline{w'\theta'})_0 / (w_* h_{\text{bl}})$.

Richardson-Number Mixing

Interior shear-driven diffusivity:

$$K_v = K_{v,0} (1 + 5 \text{ Ri})^{-1}, \quad \text{Ri} = \frac{N^2}{S^2}. \quad (109)$$

Constant Mixing

Uniform K_v (default $10^{-4} \text{ m}^2 \text{ s}^{-1}$) for idealised experiments.

Prognostic TKE closures

Two one-equation turbulent-kinetic-energy closures carry e as ocean state. The Gaspar–Burchard form (`tke`) solves

$$\frac{\partial e}{\partial t} = K_M \left[\left(\frac{\partial u}{\partial z} \right)^2 + \left(\frac{\partial v}{\partial z} \right)^2 \right] - K_H N^2 - \varepsilon + \frac{\partial}{\partial z} \left(\alpha_{\text{tke}} K_M \frac{\partial e}{\partial z} \right) \quad (110)$$

at cell interfaces and derives K_M and K_H from e and a mixing length; its options (mixing-length choice, surface boundary condition, background floors) reproduce the NEMO `zdf tke` card used by the OMIP configurations. CATKE (`catke`) follows the Oceananigans implementation, with a dynamic convective mixing length and separate Richardson-number stability functions for momentum, tracers and TKE.

Ocean Lateral Mixing

Harmonic (Laplacian) Diffusion

$F_h = A_h \nabla^2 \phi$ with constant or Smagorinsky-type coefficient.

Biharmonic (∇^4) Diffusion

$F_h = -A_4 \nabla^4 \phi$, providing scale-selective damping of grid-scale noise.

Gent–McWilliams / Redi (GM/Redi)

Mesoscale eddy parameterisation (Gent and McWilliams, 1990) with bolus velocity $\mathbf{v}^* = -\kappa_{\text{GM}} \nabla_h s$ where s is the isopycnal slope; the adiabatic stirring by unresolved eddies is represented as an advective flattening of isopycnals. Isopycnal diffusion rotates the diffusion tensor to align with isopycnal surfaces (Redi, 1982). Default $\kappa_{\text{GM}} = 1000 \text{ m}^2 \text{ s}^{-1}$.

Ocean Bottom Drag

Two formulations:

- **Linear:** $\tau_b = -r \mathbf{u}_b$ with configurable r .
- **Quadratic:** $\tau_b = -C_d |\mathbf{u}_b| \mathbf{u}_b$ with $C_d = 2.5 \times 10^{-3}$.

Ocean Surface Forcing

Three modes, available on cubed-sphere, lat-lon, and MPAS grids:

- **Prescribed:** Fixed wind stress and heat flux profiles (cosine-latitude, single-gyre, double-gyre, constant).
- **Restoring:** SST/SSS restoring toward climatology (Haney-style) with configurable timescale.
- **Bulk formulas:** COARE-like air-sea fluxes with drag, Stanton, and Dalton coefficients.

On MPAS Voronoi meshes, cell-centred wind stress (τ_x, τ_y) is projected onto TRiSK edge normals via $\tau_n = \tau_x \cos \alpha_e + \tau_y \sin \alpha_e$ where α_e is the edge normal angle. An MPAS-specific physics factory (`make_mpas_ocean_physics`) handles this projection.

Shortwave Penetration

Subsurface absorption following Jerlov (1968) water type classification (types I, IA, IB, II, III) with two-band exponential decay:

$$I(z) = I_0 [R e^{-z/\zeta_1} + (1 - R) e^{-z/\zeta_2}], \quad (111)$$

where R, ζ_1, ζ_2 depend on the water type.

Ocean Convection

Enhanced Diffusion

Convective adjustment via enhanced diffusivity K_{conv} where $N^2 < 0$, with smooth transitions.

Plume Model

Entraining mass-flux ocean convection with entrainment rate $\alpha_{\text{plume}} = 0.1$ and background detrainment $\varepsilon = 10^{-3} \text{ m}^{-1}$.

Ocean Conservation Fixers

Grid-specific volume, heat, and salt conservation fixers for cubed-sphere, lat-lon, and MPAS grids. Uniform additive correction per unit volume after each timestep. All three grids compute corrections simultaneously from the original pre-fix state (avoiding order-dependent bias) and upcast to accumulation precision before global reductions to prevent catastrophic cancellation in float32.

Limitation: These are uniform additive fixers that violate local conservation and act as spurious globally-uniform diapycnal mixing. The proper solution is flux-form tracer advection with barotropic-baroclinic flux reconciliation (Hallberg 1997, Higdon 2005).

Tracer Advection Schemes

Tracer transport in `legoesm.ocean.advection` and `legoesm.ocean.advection_som` provide upwind, TVD (Van Leer / Sweby), DST3, PPM, PPM-FCT, WENO5/7, and Second-Order Moments (SOM) schemes (Prather, 1986). The volume-weighted variance $\text{Var}(T)$ on the Eady uniform channel ranks them: `som + Smagorinsky` ($C_{\text{Smag}}=0.2$) preserves 97.1% of $\text{Var}(T)$ at 600 d in the weak-forcing regime; `tvd` preserves 94% at $\leq 2\%$ runtime overhead vs. `upwind` (recommended production default); `ppm_fct` is gated on the sign-split Zalesak/Kuzmin limiter. Forcing regime (strong $U=0.8 \text{ m s}^{-1}$ vs. weak $U=0.2$) dominates absolute mixing; the weak-forcing setup is the scientifically defensible ocean-interior benchmark.

Realistic-Geometry Spinup

Global Wolfe–Cessi spinups on real ETOPO/GEBCO bathymetry are supported on the lat-lon C-grid via `legoesm.ocean.bathymetry` (auto-detect ETOPO/GEBCO/ERDDAP/CF, Laplacian smoothing, MEO r -cap, polar-cap mask, isolated-basin filling, strait enforcement). Stability conditioners: AL81 12-corner-triad PV-flux closure with partial-cell extension (Le Sommer et al., 2009, `pv_flux_al81_partial_cell`); MOM6-style $\cos^4 \phi$ biharmonic-viscosity scaling (Griffies and Hallberg, 2000, `biharmonic_scaling_factor`); SMC03 pressure-gradient force; implicit Crank–Nicolson barotropic solver. Validated 20-yr global-overturning spinup at 1° delivers canonical thermohaline structure.

Time Integration

SSP-RK3 (Shu–Osher 1988)

The three-stage strong-stability-preserving Runge–Kutta method (Shu and Osher, 1988), whose convex-combination structure guarantees that the total-variation-diminishing property of the forward-Euler step is inherited at higher order:

$$\mathbf{u}^{(1)} = \mathbf{u}^n + \Delta t F(\mathbf{u}^n), \quad (112)$$

$$\mathbf{u}^{(2)} = \frac{3}{4}\mathbf{u}^n + \frac{1}{4}\mathbf{u}^{(1)} + \frac{1}{4}\Delta t F(\mathbf{u}^{(1)}), \quad (113)$$

$$\mathbf{u}^{n+1} = \frac{1}{3}\mathbf{u}^n + \frac{2}{3}\mathbf{u}^{(2)} + \frac{2}{3}\Delta t F(\mathbf{u}^{(2)}). \quad (114)$$

SSP-RK(4,3) and SSP-RK(5,4) variants are also available.

Robert–Asselin–Williams (RAW) Time Filter (v3.7)

The leapfrog scheme used by the semi-implicit spectral PE solver requires a time filter to damp the computational mode. The Williams (2009) modification of the Robert–Asselin filter splits the correction between the current and next time levels, restoring second-order accuracy:

$$d^n = \frac{\gamma}{2}(X^{n-1} - 2X^n + X^{n+1}), \quad (115)$$

$$X_{\text{filt}}^n = X^n + (1 - \alpha) d^n, \quad (116)$$

$$X_{\text{filt}}^{n+1} = X^{n+1} + \alpha d^n. \quad (117)$$

With $\alpha = 0.5$ (default), this is the RAW filter; $\alpha = 0$ recovers the original Robert–Asselin filter. The original RA filter introduces a first-order phase error that causes slow energy drift over climate-length integrations; the RAW modification eliminates this while preserving the same damping rate for the computational mode.

Adaptive Hyperdiffusion (v3.7)

The maximum stable coefficient for n -th order hyperdiffusion under forward Euler is

$$\nu_n = S \cdot \frac{C(n) \Delta x^n}{\Delta t}, \quad C(2) = \frac{1}{2}, C(4) = \frac{1}{8}, C(6) = \frac{1}{48}, \quad (118)$$

where $S \in (0, 1]$ is a safety factor (default 0.5). The function `adaptive_hyperdiff_coeff(dx_min, dt, order, safety)` eliminates manual tuning of the hyperdiffusion coefficient for each resolution and timestep.

Split-Explicit Methods

Each outer RK stage computes slow tendencies once, updates u, v and tracers, then runs N_s forward-backward acoustic substeps on the fast variables (w, θ', ρ') with substep size $\Delta t_s = \Delta t / N_s$.

Parallel Runtime and Precision

Every run executes on one `ParallelRuntime` (serial, single-process multi-device SPMD, multi-process MPI, or hybrid). Reverse-mode differentiation is supported across all of them, with one limitation: halo exchanges go through the `sendrecv` custom VJP, and global sums through `allreduce(SUM)`, whose identity VJP is exact only when the sum is the final loss. An intermediate global sum that is broadcast back and reused (fixer factors, normalisations, inner products) needs `broadcast_allreduce_sum`; most such call sites have not yet been migrated, so their cross-rank cotangents are dropped. MAX/MIN reductions are reserved for diagnostics.

Process and device setup

Each process is bound to exactly one GPU by narrowing `CUDA_VISIBLE_DEVICES` before the XLA backend and, on CUDA-aware MPI stacks, before `MPI_Init` (`pin_local_gpu`); without this, every rank on a node silently shares the first GPU. For multi-process CPU runs the gloo collectives are pinned to the fastest addressed network interface rather than the interface the hostname resolves to (`LEGOESM_GLOO_IFACE`; `default` restores JAX's choice). The GPU command-buffer default captures fusions and custom calls but not collectives, which measured faster on multi-rank steps; `LEGOESM_XLA_CMDBUF_COLLECTIVES=1` restores collective capture. Cluster launchers additionally set a per-lane NCCL channel recipe.

Halo schedules on the Voronoi (MPAS) mesh

Before sharding, the mesh is reordered so that each device owns one contiguous block of cells; on the SPMD path the default partition cuts a Hilbert space-filling curve into balanced chunks (recursive coordinate bisection and METIS remain selectable). On the atmosphere, edges and vertices inside each owner block are also placed in Hilbert order (`edge_order="hilbert"`, the default) so that stencil gathers touch nearby memory; the ocean lanes keep owner order. Both layouts evaluate the same discrete operators; only the memory layout differs. The halo fill is a sequence of coloured `ppermute` rounds. Because every pair in a round ships the round's largest payload, the colouring groups pairs of similar size (size-aware colouring, on by default, `LEGOESM_MPAS_SIZE_COLORING=0` to disable); this changes only the wire grouping, and results are bit-identical to the legacy colouring. Three further strategies are opt-in: a ragged all-to-all fill (`LEGOESM_MPAS_RAGGED_HALO`, with an `auto` setting that uses it only up to 32 devices, beyond which its per-slice cost loses), a wide halo that fills once per model step at the depth of all its tendency evaluations (`LEGOESM_MPAS_WIDE_HALO=1`), and an interior/rim overlap that computes the device-interior tendency while the halo is in flight (`LEGOESM_MPAS_HALO_OVERLAP=1`). Mutually exclusive strategies are refused rather than silently combined.

Lat-lon and cubed-sphere exchanges

On the lat-lon band decomposition the per-stage halo pads can be packed into a single exchange at stage entry (`LEGOESM_LATLON_PACKED_EXCHANGE=1`, opt-in in the library and enabled by the scaling launchers), and the static band-geometry pads are precomputed once instead of being exchanged every step (default on). The cubed-sphere lane keeps the native 4-D halo exchange described with the grid (all vertical levels in one message per neighbour).

Ocean barotropic solver at scale

The distributed MPAS ocean solves the implicit free surface with a fixed-iteration preconditioned conjugate gradient (default 15 iterations) with the standard two-reduction recurrence, the configuration the default was measured with; the single-reduction (Chronopoulos–Gear) form, which batches the two inner products of each iteration into one allreduce, is selectable. The default preconditioner (`barotropic_implicit_pcg_precond="gpoly"`) is a four-sweep Neumann-series polynomial evaluated on the global operator using the existing halo, so the answer does not depend on the number of devices; it replaced the device-local polynomial at 20 iterations ("`poly`"), which needs no halo exchange and remains selectable, and is the one the MPI-per-rank lane must use. A single device solves to tolerance with stock conjugate gradient whatever the preconditioner. The MPAS ocean runs multi-device through `run_omip.py --enable-mpas-spm`.

Precision policies

Precision is a policy rather than a global switch. A `PrecisionPolicy` assigns separate dtypes to storage, compute, accumulation (global sums and budget integrals) and control (implicit solves, EOS, diagnostics needing tight tolerances). Four modes are provided: `fp32`, `fp64`, `mixed` (float32 storage and compute, float64 accumulation and control) and `mixed_fp64_storage` (float64 storage, float32 compute). In mixed mode the stiff semi-implicit Helmholtz solve runs in float64 and returns the caller's dtype, surface pressure is kept in float64 as the conservation field, and sharded runs that would silently carry it in float32 raise instead. A per-mode leaf-dtype test pins which state leaves carry which dtype. The production AMIP configuration runs in `fp64`.

Land, Ice, and Coupler

Slab Land Model

Energy balance

$$C_s d_s \frac{dT_s}{dt} = (1 - \alpha_l) SW^\downarrow + \varepsilon_l (LW^\downarrow - \sigma_{SB} T_s^4) - SH - LH. \quad (119)$$

Bucket hydrology

$$\frac{dW}{dt} = P - E, \quad W \in [0, W_{\max}]. \quad (120)$$

Moisture availability $\beta = \beta_{\min} + (1 - \beta_{\min}) W/W_{\max}$ modulates evaporation via $q_{\text{sfc}} = \beta q_s(T_s)$.

Surface humidity consistency. All slab surface tiles (land, sea ice, lake) recompute q_{sfc} from the *updated* surface temperature T_s^{n+1} (and updated W^{n+1} for land) before constructing the `TileResponse`. This ensures that the returned q_{sfc} and T_{surface} are thermodynamically consistent at the end of the timestep.

Sea-Ice Model

Surface temperature

$$\rho_i c_i \frac{h_{\text{eff}}}{2} \frac{dT_i}{dt} = Q_{\text{sfc}} - \frac{k_i(T_f - T_i)}{h_{\text{eff}}}, \quad h_{\text{eff}} = \max(h_i, h_{\min}). \quad (121)$$

Growth and melt

$$\frac{dh_i}{dt} = \frac{F_{\text{cond}} - F_{\text{ocean}}}{\rho_i L_f}. \quad (122)$$

Concentration

$$\frac{dc}{dt} = \begin{cases} (dh/dt)(1-c)/h_{\text{new}} & \text{growing,} \\ (dh/dt)c/h_{\text{eff}} & \text{melting.} \end{cases} \quad (123)$$

EVP Rheology (Hunke & Dukowicz 1997)

The elastic–viscous–plastic rheology (Hunke and Dukowicz, 1997) adds an artificial elastic term to the viscous–plastic ice stress so that the momentum equation can be subcycled explicitly, avoiding an implicit solve. When `dynamics="evp"`, the ice momentum equation is

$$m \frac{\partial \mathbf{u}_i}{\partial t} = \boldsymbol{\tau}_a + \boldsymbol{\tau}_w - m f \hat{\mathbf{k}} \times \mathbf{u}_i - mg \nabla \eta + \nabla \cdot \boldsymbol{\sigma}, \quad (124)$$

where $\boldsymbol{\tau}_a, \boldsymbol{\tau}_w$ are air and ocean drag, and $\boldsymbol{\sigma}$ is the internal stress tensor from an elliptic yield curve with eccentricity $e = 2$. Ice strength:

$$P = P^* h \exp(-C(1-c)), \quad P^* = 27.5 \text{ kPa}, \quad C = 20. \quad (125)$$

The elastic term provides subcycling stability ($N_{\text{evp}} = 120$ substeps).

Multi-category ice

For $N_{\text{cat}} > 1$, per-category state (h_k, T_k, c_k) evolves independently under thermodynamics, then the linear remapping of Lipscomb (2001) redistributes ice across thickness categories.

Land Carbon Cycle and Stomatal Conductance

DALEC-990 Carbon Pools

Six prognostic pools: labile C_{lab} , foliage C_f , root C_r , wood C_w , litter C_{lit} , soil organic matter C_{SOM} . GPP via light-use efficiency; decomposition via Q_{10} temperature sensitivity with moisture limitation.

Farquhar (1980) Photosynthesis

The Farquhar et al. (1980) biochemical model expresses the leaf assimilation rate as the smooth minimum of a Rubisco-limited rate W_c and an electron-transport-limited rate W_j , less leaf respiration R_d :

$$A = \min(W_c, W_j) - R_d, \quad W_c = V_{c,\text{max}} \frac{C_i - \Gamma^*}{C_i + K_c(1 + O/K_o)}, \quad W_j = \frac{J(C_i - \Gamma^*)}{4(C_i + 2\Gamma^*)}, \quad (126)$$

where the minimum is a smooth (differentiable) function. $V_{c,\text{max}}$ and J_{max} follow peaked Arrhenius temperature responses (Bernacchi et al., 2001). Soil moisture stress: $V_{c,\text{max}} \leftarrow V_{c,\text{max}} \beta_{\text{soil}}$.

Stomatal conductance

Ball–Berry (Ball et al., 1987): $g_s = g_0 + g_1 \text{ ARH}/C_s$.

Medlyn (Medlyn et al., 2011): $g_s = g_0 + 1.6(1 + g_1/\sqrt{\text{VPD}}) A/C_s$.

Jarvis (Jarvis, 1976): $g_s = g_{s,\max} f(\text{PAR}) f(T) f(\text{VPD}) f(\theta)$ (CO₂-independent; used when carbon cycle is inactive).

The coupled A - g_s - C_i system is solved by 5 fixed-point iterations (unrolled for JAX JIT).

Plant Functional Types (PFTs)

CLM5-style weighted-average lookup over 17 Plant Functional Types (needleleaf/broadleaf trees, shrubs, grasses, crops), with 12 trainable land parameters per PFT: vegetation albedo, emissivity, roughness length, maximum soil water, soil heat capacity, soil depth, root depth, wilting point, field capacity, $V_{c,\max,25}$, LCMA, and stomatal slope g_1 . Each parameter is bounded via sigmoid transforms:

$$p = p_{\min} + (p_{\max} - p_{\min}) \sigma(x_{\text{raw}}). \quad (127)$$

Three parameter provider modes:

- **Constant:** Spatially uniform default values.
- **PFT:** Weighted average $\bar{p} = \sum_{k=1}^{17} f_k p_k$ from the PFT lookup table.
- **Neural:** 3-hidden-layer FCNN (GELU, 48 hidden units) mapping static features (lat/lon, soil type, elevation, precipitation, temperature, forest age) to land parameters.

DifferLand Training

End-to-end differentiable land model training with three carbon cycle modes: “none” (disabled), “differland” (prognostic DALEC-990 with trainable allocation and decomposition parameters), and “seasonal” (prescribed LAI). Phenology uses hemisphere-aware Gaussian windows for leaf emergence/fall. Decomposition follows Q_{10} temperature sensitivity:

$$k(T, \theta) = k_0 Q_{10}^{(T - T_{\text{ref}})/10} f_w(\theta). \quad (128)$$

NPP allocation is sequential: foliar $\rightarrow$ labile $\rightarrow$ root $\rightarrow$ wood.

Interactive Multilayer Land on the MPAS Lane

The production atmosphere couples a multilayer soil model (Richards-equation soil water, `multilayer_n_layers` = 10 layers over 3m in the production deck) with prognostic snow and a canopy energy balance selected by `land_surface_scheme` (`simple_seb`, `two_leaf` (default), or the multilayer `clm_ml` canopy, whose JAX backend is vendored in-repo). Land columns are packed so that only land cells are integrated, and the land tile may be called at its own cadence (`land_update_seconds`): the atmospheric forcing is averaged over the interval, conserving its mass and energy, and the tile’s fluxes and skin temperature are held in between. The production deck steps the land every 1800s, as CLM does under CAM6, with soil freeze/thaw enabled, calibrated land tables, seasonal refresh of leaf area, canopy height and soil albedo, and a 1979 ERA5-based soil initial state. Sea ice on this lane carries a prognostic skin temperature (`mpas_ice_skin_prognostic`).

Tile Coupler

Surface fluxes are area-weighted blends over ocean, land, sea-ice, and lake tiles:

$$X_{\text{blend}} = f_{\text{ocn}} X_{\text{ocn}} + f_{\text{land}} X_{\text{land}} + f_{\text{ice}} X_{\text{ice}} + f_{\text{lake}} X_{\text{lake}}. \quad (129)$$

Fluxes are time-averaged over each coupling window before being passed to the atmosphere.

Ocean albedo. The constant ocean albedo α_{ocn} is controlled by `CouplerConfig.ocean_albedo` (default 0.06). When `ocean_albedo_config.method="zenith"`, a Briegleb-type zenith-angle formula is used (Briegleb et al., 1986):

$$\alpha_{\text{ocn}} = \frac{0.026}{1.1 \mu_0^{1.7} + 0.065} + 0.15 (\mu_0 - 0.1)(\mu_0 - 0.5)(\mu_0 - 1.0), \quad (130)$$

where $\mu_0 = \cos \theta_z$ is the cosine of the solar zenith angle.

External forcing. Time-varying greenhouse gases (CO_2 , CH_4 , N_2O , CFCs), total and spectral solar irradiance, ozone, and tropospheric and volcanic aerosol are read from CMIP6 input4MIPs files and interpolated in time. With RRTMG/RRTMGP radiation and `ozone_forcing="external" / aerosol_forcing="external"`, ozone and aerosol optical properties enter the radiation solver; multi-year files are sampled at the simulated calendar year (`start_year`) rather than collapsed onto a 12-month cycle. Gray radiation does not respond to these channels.

Two-layer lake model. Epilimnion (mixed layer, default $h_{\text{epi}} = 5$ m) and hypolimnion (deep layer, $h_{\text{hypo}} = 20$ m) connected by wind-enhanced vertical mixing:

$$k_{\text{eff}} = k_{\text{mix}} (1 + \alpha_{\text{wind}} |\mathbf{V}|), \quad (131)$$

with surface energy balance driving the epilimnion temperature and mixing flux $F_{\text{mix}} = \rho c k_{\text{eff}} \Delta T / d_{\text{mid}}$ coupling the two layers. Lake freezing is clamped at $T_f = 273.15$ K. Bulk flux options: constant coefficients or iterative MOST.

Machine Learning: Spherical Fourier Neural Operator

The Spherical Fourier Neural Operator (Bonev et al., 2023) generalises the Fourier neural operator (Li et al., 2021) to the sphere by replacing the planar FFT with a spherical-harmonic transform, so the learned global convolution is rotation-equivariant and free of the polar artefacts a lat-lon FFT would introduce; it provides the learned-dynamics backbone in the same spirit as the neural general circulation model of Kochkov et al. (2024). The SFNO (ACE/FourCastNet v2 style) operates on the Gaussian grid.

Architecture

1. **Encoder:** Pointwise linear $n_{\text{in}} \rightarrow n_{\text{emb}}$.
2. **N processor blocks**, each containing:
 - Layer normalisation,
 - SH analysis $\rightarrow$ learned spectral convolution $\rightarrow$ SH synthesis,
 - GELU activation,
 - Pointwise MLP (two layers with expansion factor),
 - Skip connection.
3. **Decoder:** Pointwise linear $n_{\text{emb}} \rightarrow n_{\text{out}}$.
4. **Big-skip residual:** $\text{output} = \text{input} + \text{decoder}(\dots)$.

Spectral convolution

In the SH coefficient space:

$$\hat{f}_{nm}^{\text{out}} = W_{nm} \hat{f}_{nm}^{\text{in}}, \quad (132)$$

where $W_{nm} \in \mathbb{R}^{n_{\text{emb}} \times n_{\text{emb}}}$ is a learnable weight matrix per wavenumber pair.

Neural and Learned Physics

NeuralPhysics

Equinox MLP column model with hybrid mode blending learned corrections with conventional physics via a configurable blending coefficient.

LearnedColumn

Per-column physics replacement following Rasp et al. (2018), mapping the full column state to temperature and moisture tendencies via a neural network, trainable end-to-end with the dynamical core.

Training and Calibration

The segment-based training lanes (WeatherBench and AIMIP) differentiate the model itself: a segment is built once by `build_segment_fn` and its non-donating, non-jitted form is called inside `eqx.filter_value_and_grad`, with per-iteration forcing passed as an explicit argument rather than captured in a closure so that it never triggers recompilation.

One model registry

The WeatherBench and AIMIP campaigns train the same four variants — `classical` (the physics suite with trainable parameters), `column_nn` (a learned column model), `sfno_physics` and `sfno_full` — built by a single `build_variant` with one set of architecture defaults, so that “the same variant” means the same pytree in both campaigns. The classical WeatherBench base-line runs the production CAM6 scheme set on the spectral core (`config/wb/campaign/spectral_t63_cam6.yaml`) a sibling deck prescribes ERA5 surface fluxes as the lower boundary condition.

Parameter specifications and tiers

Every physics scheme configuration declares a machine-readable `__param_spec__`: units, bounds, a tunable tier, a constraining transform (sigmoid, softplus or none), a category and a literature reference for each tunable float field; float fields that are not trainable (for example prescribed gas concentrations) are listed as excluded with a reason. `build_trainable_params(tier=...)` collects tier 1 (`"core"`: well-posed parameters already trained in practice), tier ≤ 2 (`"extended"`: closure knobs with literature bounds) or tier ≤ 3 (`"aggressive"`); tier 0 marks numerical floors, smoothing widths and iteration-coupled constants that are never trainable and are refused on the command-line calibration route. Trained values are injected through the configuration pytree by `apply_param_overrides` inside the loss, so the leaves are traced during training while production keeps static values. The MIP drivers accept a YAML calibration file of tuned values keyed by the qualified name `scheme_key.field` (`--params`), validated against the declared bounds, and a CI audit checks that every specified tunable is reachable from a driver.

No inert parameters

A parameter that the forward model never consumes would appear calibrated while doing nothing. The training driver, the spectral neural-GCM trainer and the offline land trainers therefore evaluate the full-data gradient at initialisation and abort if any trainable leaf is identically zero or entirely non-finite (`assert_no_inert`); parameters belonging to inactive schemes are frozen out of the trainable set before this check.

Process-level calibration

Turbulence closures are tuned in the single-column model against matched large-eddy simulations of several boundary-layer regimes, and convection against a cloud-resolving radiative-convective equilibrium reference (`scm_rce_metrics` supplies the normalised profile score). An offline compare-to-reanalysis loop runs the model, ranks the columns with the largest bias against ERA5, diagnoses a closure coefficient from targeted LES, and keeps only rounds that reduce the bias. Optimisers come from `create_optimizer` (warm-up, cosine decay, gradient clipping; AdamW by default, Muon available).

Conservation Corrections

Atmosphere

Dry-air mass is corrected by a uniform adjustment to p_s that restores the area-weighted global integral:

$$\Delta\bar{p} = \frac{\sum_i A_i p_{s,i}^{\text{new}} - \sum_i A_i p_{s,i}^{\text{old}}}{\sum_i A_i}, \quad p_{s,i}^{\text{corr}} = p_{s,i}^{\text{new}} - \Delta\bar{p}, \quad (133)$$

so that $\sum_i A_i p_{s,i}^{\text{corr}} = \sum_i A_i p_{s,i}^{\text{old}}$ (global dry mass conserved exactly). Column moisture is rescaled: $q^{\text{corr}} = q^{\text{new}} \cdot Q_{\text{old}}/Q_{\text{new}}$.

Initial-mass anchoring (v3.7). For long AMIP/CMIP integrations, the fixer can anchor to the *initial* mass $M_0 = \sum_i A_i p_{s,i}^{\text{init}}$ rather than the previous step's mass, preventing cumulative $O(\varepsilon)$ drift over millions of timesteps. The correction becomes

$$\Delta\bar{p} = \frac{\sum_i A_i p_{s,i}^{\text{new}} - M_0}{\sum_i A_i}. \quad (134)$$

Set via `model.set_initial_mass(state)` or `anchor_mass_to_initial=True` in the config. All accumulations use float64 regardless of the state precision.

Water and Energy on the Production Lane

Three corrections close the water and energy budgets of the hybrid-coordinate MPAS lane used for AMIP.

- **Conservative tracer vertical transport.** All tracers on the hybrid coordinate are advected vertically by a van Leer scheme written in increment (face-minus-cell) form with the interfacial mass flux and each column's layer pressure thickness, so the discrete transport satisfies the product rule with continuity. The earlier advective up-wind operator on averaged mass fluxes created water; there is no option to restore it on this lane. Potential temperature uses the conservative Simmons–Burridge flux form (`mpas_vert_advection_scheme="sb"` in the production deck).
- **Column-conserving positivity.** Negative tracer values are set to zero and the column's positive values rescaled so that the layer-mass-weighted column integral is unchanged, instead of a plain clip that creates mass (`mpas_conservative_tracer_clamp=True`, the default, applied on every grid despite its name). Each per-mass tracer, including the per-mass number concentrations of the two-moment schemes, is conserved separately.
- **Energy-consistent moisture floor.** When the vapour floor removes moisture, the latent heat of the removed vapour is taken out of the temperature equation as well (`energy_consistent_moisture_clip`; off in the library, on in the production deck).

Global water and energy residuals (`diagnostics.water_budget`, `diagnostics.energy_budget`) are diagnostic tripwires, never fed back into the physics: they use the coordinate’s hybrid-aware layer mass, area-weighted means, and AD-safe global sums, so the reported residual does not depend on how the run is sharded beyond floating-point round-off.

Ocean

Volume conservation: $\eta^{\text{corr}} = \eta^{\text{new}} + (\sum \eta_{\text{old}} A - \sum \eta_{\text{new}} A) / \sum A$.

Heat and salt conservation: uniform additive correction per unit volume.

Physical Constants

All physical constants are defined once in `legoesm.constants`; a CI ratchet refuses new hard-coded copies elsewhere (the FV3 duo-grid lane deliberately uses FV3’s own constants).

Symbol	Name	Value	Unit
g	Gravitational acceleration	9.80616	m s^{-2}
R_d	Dry-air gas constant	287.05	$\text{J kg}^{-1} \text{K}^{-1}$
c_p	Dry-air heat capacity	1004.64	$\text{J kg}^{-1} \text{K}^{-1}$
R_v	Water-vapour gas constant	461.51	$\text{J kg}^{-1} \text{K}^{-1}$
c_{pv}	Water-vapour heat capacity	1846.0	$\text{J kg}^{-1} \text{K}^{-1}$
c_{pw}	Liquid-water heat capacity	4218.0	$\text{J kg}^{-1} \text{K}^{-1}$
c_{pi}	Ice heat capacity	2106.0	$\text{J kg}^{-1} \text{K}^{-1}$
L_v	Latent heat of vapourisation at 0°C	2.501×10^6	J kg^{-1}
L_f	Latent heat of fusion at 0°C	3.337×10^5	J kg^{-1}
L_s	Latent heat of sublimation at 0°C , $L_v + L_f$	2.8347×10^6	J kg^{-1}
κ	R_d/c_p	0.2857	—
p_0	Reference pressure	10^5	Pa
σ_{SB}	Stefan–Boltzmann constant	5.670374×10^{-8}	$\text{W m}^{-2} \text{K}^{-4}$
κ_{vK}	von Kármán constant	0.4	—
Ω	Earth rotation rate	7.292×10^{-5}	rad s^{-1}
$R_\oplus$	Earth radius	6.371229×10^6	m
ρ_0	Ocean reference density	1025	kg m^{-3}

One latent-heat family. Latent heats are temperature dependent and share one Kirchhoff family in `legoesm.thermo`:

$$L_v(T) = L_v^0 - (c_{pw} - c_{pv})(T - T_0), \quad (135)$$

$$L_s(T) = L_s^0 + (c_{pv} - c_{pi})(T - T_0), \quad (136)$$

$$L_f(T) = L_s(T) - L_v(T), \quad (137)$$

with L_v^0 , L_s^0 the values at $T_0 = 273.15 \text{ K}$, so the three latent heats close identically at every temperature (L_v falls by 2372 J kg^{-1} per kelvin, about 3% lower at 30°C). The latent heat of water leaving a surface is the phase-weighted blend $(1-f)L_v(T_{\text{sfc}}) + fL_s(T_{\text{sfc}})$, with f the frozen fraction of the evaporating surface, so the flux stays continuous and differentiable across melt-out; a surface made of several evaporating components (canopy under snow, sea-ice categories) sums $L E$ per component. Remaining uses of the bare 0°C constants at component interfaces are tracked by a shrink-only ratchet.

Ocean Biogeochemistry

Two schemes: abiotic (DIC + ALK only) and NPZD ecosystem.

Carbonate Chemistry and Air-Sea CO₂ Exchange

The ocean surface $p\text{CO}_2$ is computed from DIC, ALK, T , S via the carbonate-system solution of Follows et al. (2006). Air-sea flux (Wanninkhof, 2014):

$$F_{\text{CO}_2} = k_w K_0 (p\text{CO}_{2,\text{ocean}} - p\text{CO}_{2,\text{atm}}), \quad (138)$$

where $k_w = 0.251 U_{10}^2 (\text{Sc}/660)^{-1/2}$ is the gas transfer velocity [cm/hr] and K_0 is Henry's law solubility.

NPZD Ecosystem Model

Four biological tracers: NO_3 , phytoplankton P , zooplankton Z , detritus D (Fasham et al., 1990).

$$\frac{\partial P}{\partial t} = \mu_P P - G - m_P P, \quad (139)$$

$$\frac{\partial Z}{\partial t} = \gamma_Z G - m_Z Z^2, \quad (140)$$

$$\frac{\partial D}{\partial t} = (1 - \gamma_Z)G + m_P P + m_Z Z^2 - r_D D + w_s \frac{\partial D}{\partial z}, \quad (141)$$

$$\frac{\partial \text{NO}_3}{\partial t} = -\mu_P P + r_D D, \quad (142)$$

where $\mu_P = \mu_{\max} f(T) \min(L_N, L_I)$ with Michaelis-Menten nutrient limitation $L_N = \text{NO}_3/(k_N + \text{NO}_3)$ and Beer-Lambert light limitation with chlorophyll self-shading. Grazing G follows Holling-III. Redfield stoichiometry ($R_{C:N} = 6.625$) couples biology to DIC/ALK.

CMIP Infrastructure

CF/CMOR-Compliant Output

The `CFWriter` class produces CF-1.8 NetCDF4 files with CMIP6 Data Reference Syntax (DRS) filenames: `<var>.<table>.<model>.<expt>.<variant>.<grid>.nc`. Variable metadata (units, `cell_methods`, `positive`, `cell_measures`) is read from the official CMIP6 CMOR tables vendored with the package rather than retyped; adding an output variable means adding it to the table-subset builder, never editing metadata by hand. Eleven tables are available: `Amon` (45 variables), `Aday`, `day`, `CFday`, `fx`, `Lmon`, `Omon`, `Oyr`, `Ofx`, `SImon` and `SIyr`. 3-D fields are interpolated to the standard 19 CMIP6 pressure levels read from the same tables.

On the MPAS lane every published time-varying field other than the daily extremes is an interval mean whose `cell_methods` is true: radiative and surface fluxes are accumulated every step, and state fields (near-surface and level temperature, winds, humidity, cloud cover, water paths) are accumulated from hourly samples so that monthly means are not aliased onto one hour of the diurnal cycle. Daily maxima and minima are taken from the same hourly samples. Multi-rank MPAS runs gather the cells each rank owns before writing, and a multi-rank configuration that would write empty files is refused. Across chained restarts the writer keeps the first value of a time already on disk instead of writing a duplicate, and refuses any time earlier than the last written one, so the time axis stays monotonic.

Production AMIP Configuration

The production AMIP configuration is a single machine-independent file, `config/amip/amip_production.yaml`, run with `scripts/run/run_amip.py --config`; machine-specific paths are supplied separately by `config/amip/amip_production.sh`. It homogenises the atmosphere with CAM6:

- **Grid and dynamics:** MPAS Voronoi mesh at resolution 6 (40,962 cells, $\sim 1.1^\circ$), the CAM6 32-level hybrid table (model top 2.255 hPa), dynamics time step 112.5 s, float64.
- **Cadence:** physics every 1800 s (16 dynamics steps), radiation hourly, and three cloud macro/microphysics sub-steps per physics step, as in CAM6.
- **Physics:** RRTMG-band radiation with McICA cloud overlap; Zhang–McFarlane deep convection; prognostic CLUBB turbulence with CAM6 cloud fraction; two-moment Morrison microphysics in its MG2 flavour; McFarlane orographic gravity-wave drag; CESM Large–Yeager surface fluxes; homogeneous ice nucleation and aerosol-dependent cloud droplet number.
- **Surface:** interactive multilayer land with a two-leaf canopy (Section 15.2 and the land section above) and prognostic sea-ice skin temperature over prescribed SST/sea-ice concentration.
- **Forcing:** CMIP6 input4MIPs SST/SIC, greenhouse gases, ozone, aerosol and spectral solar irradiance, starting 1979-01-01 from ERA5.

The deck names its scientific choices explicitly, including many values that equal a code default, so that the deck rather than the library defaults is the record of the run; a few choices are still inherited from code defaults (for example the Morrison flavour, `mg`). The previous Sundqvist/L36 production deck is preserved as `config/amip/amip_sundqvist_l36.yaml`.

The CMIP6 forcing channels are read by loaders that all carry `start_year`, so simulation days map onto each file’s absolute time axis: SST/SIC (monthly boundary conditions), annual greenhouse gases ($\text{CO}_2/\text{CH}_4/\text{N}_2\text{O}/\text{CFC-11}/\text{CFC-12}$; consumed by RRTMG/RRTMGP only — gray radiation does not respond), ozone (monthly climatology or multi-decadal transient), solar TSI plus spectral band fractions, tropospheric aerosol, and volcanic aerosol sampled at the correct calendar months. `scripts/data/generate_amip_forcing.py` produces synthetic CMIP6-shaped replacements when the files are absent, and `scripts/validate/smoke_test_amip_all_grids.py` exercises the dispatch path on the cubed-sphere, lat-lon, Gaussian spectral and Voronoi/MPAS grids with gray radiation; it does not validate the full forcing stack on every grid.

Experiment Templates

Six CMIP6-style experiments with built-in GHG time series:

Template	Period	Forcing	CO ₂ range
<code>piControl</code>	1850– ∞	Fixed 1850	284.3 ppmv
<code>historical</code>	1850–2014	Transient	284→397 ppmv
<code>ssp245</code>	2015–2100	SSP2-4.5	397→600 ppmv
<code>ssp585</code>	2015–2100	SSP5-8.5	397→1135 ppmv
<code>amip</code>	1979–2014	Prescribed SST	Transient
<code>1pctCO2</code>	1–150	1%/yr increase	284→1277 ppmv

Restart and Reproducibility

Each restart file is accompanied by a `.meta.json` sidecar containing: SHA-256 state digest (sorted, concatenated array bytes), SHA-256 config hash, platform identifier, JAX version, x64 mode, and git hash. The function `verify_reproducibility()` compares two restart files and reports any discrepancies.

Tuning Guide

A registry of 16 tunable parameters across five categories (dynamics, radiation, convection, diffusion, surface) with valid ranges, default values, sensitivity ratings, and physical notes. `validate_tuning(config)` checks CFL consistency, range violations, and conflicting options. `recommended_params(resolution, nlev)` provides resolution-appropriate defaults.

OMIP Simulation Pipeline

The `run_omip.py` script provides a complete Ocean Model Intercomparison Project (OMIP) simulation pipeline supporting cubed-sphere, lat-lon, MPAS, and spectral grids. Features include WOA18-based initialisation with fallback, Haney-style SST/SSS restoring toward climatology, configurable physics presets (“none”, “minimal”, “full” with KPP, GM/Redi, bottom drag, shortwave penetration), per-grid resolution and timestep defaults, and diagnostic extraction with wall-time tracking. The lat-lon and Voronoi lanes run multi-device (`--enable-latlon-spm`, `--enable-mpas-spm`).

Reference-model fidelity. Ocean numerics are checked term by term against reference models through recipes that select the shared canonical blocks rather than bespoke solvers: NEMO (the GYRE and DINO configurations and NEMO’s idealised test cases), MITgcm, Oceananigans and Veros. Glue that exists only to mimic a reference (halo stripping, axis transposes, time-level bookkeeping) lives in the fidelity harness, never in the model, and oracle tiers are locked whenever a truth tier (conservation, equivariance, analytic solution) fails.

4D-Var Data Assimilation (v3.9)

The model includes a fully differentiable 4D-Var data assimilation system (Courtier et al., 1994; Rabier et al., 2000) that exploits JAX’s automatic differentiation to compute exact gradients of the cost function without hand-coded adjoints, replacing the conventional separately-maintained tangent-linear and adjoint models.

Cost Function

The 4D-Var cost function minimised over the assimilation window $[t_0, t_N]$ is

$$J(\mathbf{x}_0) = \frac{1}{2}(\mathbf{x}_0 - \mathbf{x}_b)^T \mathbf{B}^{-1}(\mathbf{x}_0 - \mathbf{x}_b) + \frac{1}{2} \sum_{i=1}^N (\mathcal{H}_i[\mathcal{M}_{0 \rightarrow i}(\mathbf{x}_0)] - \mathbf{y}_i)^T \mathbf{R}_i^{-1}(\mathcal{H}_i[\mathcal{M}_{0 \rightarrow i}(\mathbf{x}_0)] - \mathbf{y}_i), \quad (143)$$

where $\mathbf{x}_b$ is the background state, $\mathbf{B}$ the background error covariance, $\mathcal{M}_{0 \rightarrow i}$ the nonlinear model propagator, $\mathcal{H}_i$ the observation operator at time t_i , $\mathbf{y}_i$ the observations, and $\mathbf{R}_i$ the observation error covariance. Gradients $\nabla_{\mathbf{x}_0} J$ are computed via `jax.grad` through the full forward model trajectory.

Control Vector

The control vector comprises perturbations to:

- Horizontal wind components (u , v),
- Temperature (T),
- Specific humidity (q_v),
- Surface pressure (p_s).

Increments are applied in model space and can optionally be preconditioned by the square root of $\mathbf{B}$ (i.e. $\mathbf{x}_0 = \mathbf{x}_b + \mathbf{B}^{1/2} \boldsymbol{\chi}$ with minimisation in $\boldsymbol{\chi}$ -space).

Background Error Covariance

Four background error covariance formulations are implemented:

1. **Diagonal:** $\mathbf{B} = \text{diag}(\sigma_b^2)$, the simplest variance-only model.
2. **Diffusion:** $\mathbf{B} = \mathbf{\Sigma} \mathbf{C} \mathbf{\Sigma}$ where $\mathbf{C}$ is a correlation matrix defined implicitly via iterated Laplacian diffusion on the model grid.
3. **Spectral:** $\mathbf{B}$ is diagonal in spectral space, with variance specified per total wavenumber; correlations arise naturally from the spectral truncation.
4. **Hybrid:** $\mathbf{B} = \beta_1 \mathbf{B}_{\text{static}} + \beta_2 \mathbf{B}_{\text{ensemble}}$, blending a static covariance with flow-dependent ensemble covariance, the latter spatially localised by a compactly supported Gaspari and Cohn (1999) correlation function to suppress spurious long-range sampling error.

A GEN_BE-style operator is also available, $\mathbf{B}^{1/2} = U_{\text{wind}} U_{\sigma} U_{\text{vert}} U_{\text{bal}} U_{\text{horiz}}$, whose horizontal correlation, balance regression, vertical EOFs and variances are fitted from an ensemble of background-error samples (for example NMC forecast differences); its horizontal correlation uses spherical harmonics on Gaussian grids and neighbour-graph diffusion on Voronoi meshes.

Observation Operators

Three observation operator types are provided:

- **Direct:** identity mapping from model grid points to observed quantities.
- **Interpolating:** bilinear (horizontal) and linear (vertical) interpolation from the model grid to observation locations.
- **Composite:** chains multiple operators (e.g. interpolation followed by a forward radiative transfer calculation).

Minimiser

Two minimisers run entirely on device inside `jax.lax.while_loop`: limited-memory BFGS with a backtracking line search, and preconditioned Polak–Ribière conjugate gradient. The preconditioner is the square root of $\mathbf{B}$, transforming the problem to χ -space where the background term becomes isotropic. An incremental 4D-Var driver wraps them in an outer relinearisation loop. Both support configurable iteration counts and convergence tolerances.

DA Cycling and Incremental Analysis

For operational-style cycling, the analysis from each window serves as the background for the next. An incremental analysis update (IAU) option distributes the analysis increment gradually over the subsequent forecast window rather than inserting it as an instantaneous shock, reducing initialisation imbalance and gravity-wave noise.

References

References

- Audusse, E., F. Bouchut, M.-O. Bristeau, R. Klein, and B. Perthame, 2004: A fast and stable well-balanced scheme with hydrostatic reconstruction for shallow water flows. *SIAM J. Sci. Comput.*, **25**, 2050–2065.
- Ball, J.T., I.E. Woodrow, and J.A. Berry, 1987: A model predicting stomatal conductance and its contribution to the control of photosynthesis under different environmental conditions. *Prog. in Photosynthesis Res.*, **4**, 221–224.

- Bernacchi, C.J., E.L. Singsaas, C. Pimentel, A.R. Portis, and S.P. Long, 2001: Improved temperature response functions for models of Rubisco-limited photosynthesis. *Plant Cell Environ.*, **24**, 253–259.
- Bourke, W., 1972: An efficient, one-level, primitive-equation spectral model. *Mon. Wea. Rev.*, **100**, 683–689.
- Businger, J.A., J.C. Wyngaard, Y. Izumi, and E.F. Bradley, 1971: Flux-profile relationships in the atmospheric surface layer. *J. Atmos. Sci.*, **28**, 181–189.
- Colella, P. and P.R. Woodward, 1984: The Piecewise Parabolic Method (PPM) for gas-dynamical simulations. *J. Comput. Phys.*, **54**, 174–201.
- Fairall, C.W., E.F. Bradley, J.E. Hare, A.A. Grachev, and J.B. Edson, 2003: Bulk parameterization of air–sea fluxes: Updates and verification for the COARE algorithm. *J. Climate*, **16**, 571–591.
- Farquhar, G.D., S. von Caemmerer, and J.A. Berry, 1980: A biochemical model of photosynthetic CO₂ assimilation in leaves of C3 species. *Planta*, **149**, 78–90.
- Fasham, M.J.R., H.W. Ducklow, and S.M. McKelvie, 1990: A nitrogen-based model of plankton dynamics in the oceanic mixed layer. *J. Mar. Res.*, **48**, 591–639.
- Frierson, D.M.W., I.M. Held, and P. Zurita-Gotor, 2006: A gray-radiation aquaplanet moist GCM. Part I: Static stability and eddy scale. *J. Atmos. Sci.*, **63**, 2548–2566.
- Frierson, D.M.W., 2007: The dynamics of idealized convection schemes and their effect on the zonally averaged tropical circulation. *J. Atmos. Sci.*, **64**, 1959–1976.
- Gent, P.R. and J.C. McWilliams, 1990: Isopycnal mixing in ocean circulation models. *J. Phys. Oceanogr.*, **20**, 150–155.
- Griffies, S.M. and R.W. Hallberg, 2000: Biharmonic friction with a Smagorinsky-like viscosity for use in large-scale eddy-permitting ocean models. *Mon. Wea. Rev.*, **128**, 2935–2946.
- Held, I.M. and M.J. Suarez, 1994: A proposal for the intercomparison of the dynamical cores of atmospheric general circulation models. *Bull. Amer. Meteor. Soc.*, **75**, 1825–1830.
- Prather, M.J., 1986: Numerical advection by conservation of second-order moments. *J. Geophys. Res.*, **91** (D6), 6671–6681.
- Hunke, E.C. and J.K. Dukowicz, 1997: An elastic-viscous-plastic model for sea ice dynamics. *J. Phys. Oceanogr.*, **27**, 1849–1867.
- Jiang, G.-S. and C.-W. Shu, 1996: Efficient implementation of weighted ENO schemes. *J. Comput. Phys.*, **126**, 202–228.
- Kessler, E., 1969: On the Distribution and Continuity of Water Substance in Atmospheric Circulations. *Meteor. Monogr.*, No. 32, Amer. Meteor. Soc., 84 pp.
- Le Sommer, J., T. Penduff, S. Theetten, G. Madec, and B. Barnier, 2009: How momentum advection schemes influence current–topography interactions at eddy permitting resolution. *Ocean Modelling*, **29**, 1–14.
- Large, W.G., J.C. McWilliams, and S.C. Doney, 1994: Oceanic vertical mixing: A review and a model with a nonlocal boundary layer parameterization. *Rev. Geophys.*, **32**, 363–403.

- Large, W.G. and S.G. Yeager, 2004: Diurnal to decadal global forcing for ocean and sea-ice models: The data sets and flux climatologies. *NCAR Technical Note*, NCAR/TN-460+STR.
- Lin, S.-J., 2004: A “vertically Lagrangian” finite-volume dynamical core for global models. *Mon. Wea. Rev.*, **132**, 2293–2307.
- Lipscomb, W.H., 2001: Remapping the thickness distribution in sea ice models. *J. Geophys. Res.*, **106**(C7), 13989–14000.
- Louis, J.-F., 1979: A parametric model of vertical eddy fluxes in the atmosphere. *Bound.-Layer Meteor.*, **17**, 187–202.
- Medlyn, B.E., et al., 2011: Reconciling the optimal and empirical approaches to modelling stomatal conductance. *Global Change Biol.*, **17**, 2134–2144.
- Morrison, H., J.A. Curry, and V.I. Khvorostyanov, 2005: A new double-moment microphysics parameterization for application in cloud and climate models. Part I: Description. *J. Atmos. Sci.*, **62**, 1665–1677.
- Pincus, R., E.J. Mlawer, and J.S. Delamere, 2019: Balancing accuracy, efficiency, and flexibility in radiation calculations for dynamical models. *J. Adv. Model. Earth Syst.*, **11**, 3074–3089.
- Mouallem, J., L.M. Harris, and R. Benson, 2023: A duo-grid methodology for cubed-sphere shallow water models. *J. Comput. Phys.*, **485**, 112097.
- Putman, W.M. and S.-J. Lin, 2007: Finite-volume transport on various cubed-sphere grids. *J. Comput. Phys.*, **227**, 55–78.
- Seifert, A. and K.D. Beheng, 2001: A two-moment cloud microphysics parameterization for mixed-phase clouds. *Meteor. Atmos. Phys.*, **73**, 67–88.
- Shu, C.-W. and S. Osher, 1988: Efficient implementation of essentially non-oscillatory shock-capturing schemes. *J. Comput. Phys.*, **77**, 439–471.
- Simmons, A.J. and D.M. Burridge, 1981: An energy and angular-momentum conserving vertical finite-difference scheme and hybrid vertical coordinates. *Mon. Wea. Rev.*, **109**, 758–766.
- Skamarock, W.C. and J.B. Klemp, 2008: A time-split nonhydrostatic atmospheric model for weather research and forecasting applications. *J. Comput. Phys.*, **227**, 3465–3485.
- van Leer, B., 1979: Towards the ultimate conservative difference scheme. V. A second-order sequel to Godunov’s method. *J. Comput. Phys.*, **32**, 101–136.
- Wanninkhof, R., 2014: Relationship between wind speed and gas exchange over the ocean revisited. *Limnol. Oceanogr. Methods*, **12**, 351–362.
- Williams, P.D., 2009: A proposed modification to the Robert–Asselin time filter. *Mon. Wea. Rev.*, **137**, 2538–2546.
- Williamson, D.L., J.B. Drake, J.J. Hack, R. Jakob, and P.N. Swarztrauber, 1992: A standard test set for numerical approximations to the shallow water equations in spherical geometry. *J. Comput. Phys.*, **102**, 211–224.
- Wright, D.G., 1997: An equation of state for use in ocean models: Eckart’s formula revisited. *J. Atmos. Ocean. Tech.*, **14**, 735–740.
- Bonev, B., T. Kurth, C. Hundt, J. Pathak, M. Baust, K. Kashinath, and A. Anandkumar, 2023: Spherical Fourier neural operators: Learning stable dynamics on the sphere. *Proc. 40th Int. Conf. on Machine Learning (ICML)*, PMLR 202, 2806–2823.

- Courtier, P., J.-N. Thépaut, and A. Hollingsworth, 1994: A strategy for operational implementation of 4D-Var, using an incremental approach. *Quart. J. Roy. Meteor. Soc.*, **120**, 1367–1387.
- Follows, M.J., T. Ito, and S. Dutkiewicz, 2006: On the solution of the carbonate chemistry system in ocean biogeochemistry models. *Ocean Modelling*, **12**, 290–301.
- Gaspari, G. and S.E. Cohn, 1999: Construction of correlation functions in two and three dimensions. *Quart. J. Roy. Meteor. Soc.*, **125**, 723–757.
- Kochkov, D., et al., 2024: Neural general circulation models for weather and climate. *Nature*, **632**, 1060–1066.
- Li, Z., N. Kovachki, K. Azizzadenesheli, B. Liu, K. Bhattacharya, A. Stuart, and A. Anandkumar, 2021: Fourier neural operator for parametric partial differential equations. *Proc. Int. Conf. on Learning Representations (ICLR)*.
- Lin, S.-J. and R.B. Rood, 1996: Multidimensional flux-form semi-Lagrangian transport schemes. *Mon. Wea. Rev.*, **124**, 2046–2070.
- Lindzen, R.S., 1981: Turbulence and stress owing to gravity wave and tidal breakdown. *J. Geophys. Res.*, **86**(C10), 9707–9714.
- Marshall, J., A. Adcroft, C. Hill, L. Perelman, and C. Heisey, 1997: A finite-volume, incompressible Navier–Stokes model for studies of the ocean on parallel computers. *J. Geophys. Res.*, **102**(C3), 5753–5766.
- McFarlane, N.A., 1987: The effect of orographically excited gravity wave drag on the general circulation of the lower stratosphere and troposphere. *J. Atmos. Sci.*, **44**, 1775–1800.
- Mlawer, E.J., S.J. Taubman, P.D. Brown, M.J. Iacono, and S.A. Clough, 1997: Radiative transfer for inhomogeneous atmospheres: RRTM, a validated correlated- k model for the longwave. *J. Geophys. Res.*, **102**(D14), 16663–16682.
- Rabier, F., H. Järvinen, E. Klinker, J.-F. Mahfouf, and A. Simmons, 2000: The ECMWF operational implementation of four-dimensional variational assimilation. Part I: Experimental results with simplified physics. *Quart. J. Roy. Meteor. Soc.*, **126**, 1143–1170.
- Redi, M.H., 1982: Oceanic isopycnal mixing by coordinate rotation. *J. Phys. Oceanogr.*, **12**, 1154–1158.
- Ronchi, C., R. Iacono, and P.S. Paolucci, 1996: The “cubed sphere”: A new method for the solution of partial differential equations in spherical geometry. *J. Comput. Phys.*, **124**, 93–114.
- Sadourny, R., 1972: Conservative finite-difference approximations of the primitive equations on quasi-uniform spherical grids. *Mon. Wea. Rev.*, **100**, 136–144.
- Skamarock, W.C., J.B. Klemp, M.G. Duda, L.D. Fowler, S.-H. Park, and T.D. Ringler, 2012: A multiscale nonhydrostatic atmospheric model using centroidal Voronoi tessellations and C-grid staggering. *Mon. Wea. Rev.*, **140**, 3090–3105.
- Smagorinsky, J., 1963: General circulation experiments with the primitive equations. I. The basic experiment. *Mon. Wea. Rev.*, **91**, 99–164.
- Sundqvist, H., E. Berge, and J.E. Kristjánsson, 1989: Condensation and cloud parameterization studies with a mesoscale numerical weather prediction model. *Mon. Wea. Rev.*, **117**, 1641–1657.

- Thompson, G., P.R. Field, R.M. Rasmussen, and W.D. Hall, 2008: Explicit forecasts of winter precipitation using an improved bulk microphysics scheme. Part II: Implementation of a new snow parameterization. *Mon. Wea. Rev.*, **136**, 5095–5115.
- Briegleb, B.P., P. Minnis, V. Ramanathan, and E. Harrison, 1986: Comparison of regional clear-sky albedos inferred from satellite observations and model computations. *J. Climate Appl. Meteor.*, **25**, 214–226.
- Hines, C.O., 1997: Doppler-spread parameterization of gravity-wave momentum deposition in the middle atmosphere. *J. Atmos. Solar-Terr. Phys.*, **59**, 371–400.
- Jarvis, P.G., 1976: The interpretation of the variations in leaf water potential and stomatal conductance found in canopies in the field. *Philos. Trans. Roy. Soc. London B*, **273**, 593–610.
- Kuo, H.L., 1974: Further studies of the parameterization of the influence of cumulus convection on large-scale flow. *J. Atmos. Sci.*, **31**, 1232–1240.
- Rasp, S., M.S. Pritchard, and P. Gentine, 2018: Deep learning to represent subgrid processes in climate models. *Proc. Natl. Acad. Sci. USA*, **115**, 9684–9689.
- Xu, K.-M. and D.A. Randall, 1996: A semiempirical cloudiness parameterization for use in climate models. *J. Atmos. Sci.*, **53**, 3084–3102.